



# MedEvo: Medical Reasoning Evolution through Knowledge Navigation

Wang Xiaoliang<sup>♥\*</sup>, Zhuo Chen<sup>\*,</sup>, Chuanyang Zheng<sup>\*</sup>, Quanbin Wang<sup>\*</sup>, Jiangwei Lao<sup>\*</sup>, Yingao Liu<sup>\*</sup>  
Jiajia Liu<sup>\*</sup>, Jinming Cao<sup>\*</sup>, Honglin Lin<sup>♦</sup>, Zhihong Lu<sup>\*</sup>, Hao Jin<sup>♥</sup>, Jian Wang<sup>\*</sup>, Lijun Wu<sup>♦†</sup>

<sup>♥</sup>Zhejiang University   <sup>\*</sup>Ant Group   <sup>\*</sup>Shanghai Artificial Intelligence Laboratory   <sup>♦</sup>Shanghai Jiao Tong University

## Abstract

Medical multimodal reasoning models are typically trained on **teacher-generated rationales**, discarding the valid reasoning in students' own failed trajectories, while general-domain correction methods offer little guidance for errors in image reading, anatomical context, or differential diagnosis. We propose **MedEvo**, which uses **medical knowledge**, namely a tag taxonomy and error-type diagnosis, to navigate both **trajectory repair and data selection**. For each failed rollout, a **gold-blind auditor** locates the first consequential error and provides a minimal, answer-free hint; the same student **retains the valid prefix and regenerates the erroneous suffix**. After each round, MedEvo identifies capability gaps across medical tags on a **held-out feedback set** and prioritizes related questions for new rollouts and repairs. With **SFT alone**, MedEvo raises Qwen3.5-9B mean accuracy across seven benchmarks from **65.57% to 69.12%** (47.40% to **54.50%** on MedXpertQA-MM), surpassing teacher-distillation SFT (66.93%). It also raises Qwen3.5-35B and Qwen3.5-397B mean accuracy from 67.91% to **73.04%** and from 71.83% to **74.95%**, respectively. Our analyses show that failed trajectories contain reasoning worth preserving, that **repair success does not imply learning value**, and that **independent feedback is key to sustained improvement**.

**Date:** September 2026

**Code:** <https://github.com/pork-pay/MedEvo>

**Project:** <https://medevo.site/>


## 1 Introduction

Medical multimodal reasoning requires connecting visual findings with anatomical context, disease knowledge, and clinical decisions, and expert-level questions remain challenging for current medical vision-language models [2, 3, 4]. A model may correctly recognize several findings yet reach a wrong answer after misreading one image feature, overlooking a clinical clue, or making an unsupported differential diagnosis. Such **failures are often local**: the reasoning before the mistake can remain valid, so preserving valid steps and recovering from the point of failure offers a direct way to improve medical reasoning [5, 6].

Existing approaches typically train students on teacher-generated rationales or on answer-correct student rollouts [7, 8, 9]. Discarding a failed trajectory loses its valid reasoning, while replacing it with a complete teacher

<sup>\*</sup>Equal contribution.

<sup>†</sup>Corresponding author.



**Figure 1** Average accuracy across seven medical multimodal benchmarks. The MedEvo models with 9B, 35B, and 397B parameters achieve strong performance across model scales, with MedEvo-9B outperforming several larger general-purpose and medical models in this comparison. Abbreviations: billion parameters (B), supervised fine-tuning (SFT), and Hindsight Hint Distillation (HHD) [1].

solution gives the student a path it did not produce. Methods that learn from failures, such as mistake-correction pairs [6], trajectory fusion [10], and hint-guided recovery [1], were developed for mathematics and software engineering, and offer little guidance when an error depends on medical images, anatomy, or differential diagnosis. Moreover, a model’s failures reveal which medical capabilities need further training, yet this signal is rarely used to guide the next round of supervision. We therefore ask two questions: how can we repair reasoning while **preserving valid steps**, and how can **failures guide subsequent learning**?

We introduce MedEvo through the principle of **Knowledge Navigation**: using structured medical knowledge and failure diagnosis to guide **where to intervene, how to recover, and what to learn next**. This definition follows from the two questions above: improving reasoning requires both a local recovery decision and a learning decision that reflects the student’s remaining weaknesses. Medical tags locate a failure in its clinical context, while error types identify the reasoning operation that needs correction. We call this navigation because these signals direct a path from an observed failure to a targeted intervention, at two connected levels: **within a trajectory** and **across training rounds** (Fig. 4). For a failed rollout, a gold-blind auditor identifies the first consequential error, preserves the valid prefix, and provides a **minimal, answer-free hint**; the same student then regenerates the erroneous suffix, so the correction remains grounded in its own reasoning (Fig. 2). After each round, MedEvo measures performance and errors across medical tags on a **held-out feedback set** and prioritizes questions for fresh rollouts and repairs, so that recurring failures shape subsequent training.


Across seven medical multimodal benchmarks (Fig. 1), MedEvo improves the mean accuracy of Qwen3.5-9B, Qwen3.5-35B, and Qwen3.5-397B from 65.57%, 67.91%, and 71.83% to **69.12%**, **73.04%**, and **74.95%**, respectively. With Qwen3.5-9B, it outperforms teacher-distillation SFT (69.12% vs. 66.93%); with Qwen3.5-35B and Qwen3.5-397B, it surpasses GPT-5.6-sol (72.47%); Gemini-3.8-Flash has the highest mean accuracy in this comparison (76.47%).

Our contributions are threefold:

- **Medically informed local repair.** Within trajectories, Knowledge Navigation uses a medical tag taxonomy and error-type diagnosis to locate the first consequential error and a safe rollback boundary; a gold-blind auditor provides an answer-free hint, and the same student regenerates the erroneous suffix while preserving its valid prefix.
- **Feedback-driven data evolution.** Across rounds, Knowledge Navigation uses feedback: capability gaps across

medical tags and error types on a held-out feedback set prioritize the questions and failures used for fresh rollouts and repairs.

- **Empirical gains and findings.** With SFT alone, MedEvo improves Qwen3.5-9B, Qwen3.5-35B, and Qwen3.5-397B across seven benchmarks. Our analyses show that failed trajectories contain reasoning worth preserving, that repair success does not imply learning value, and that independent feedback is key to sustained improvement.



**Figure 2** Local navigation from a failed reasoning trajectory. A student produces an incorrect answer; a teacher retains a supported prefix and supplies an answer-blind hint; the same student re-examines the evidence and continues to a corrected answer. The comic illustrates the interface; localization and admission checks are detailed in §3.

## 2 Related Work


### 2.1 Medical Multimodal Reasoning and Supervision Quality

Medical vision-language models have been built through biomedical instruction tuning [2, 3], generalist medical training in Lingshu [11], and bilingual image understanding in BiMediX2 [12]. MedGemma provides medical vision-language foundation models based on Gemma 3 [13]. Beyond answer accuracy, **process-level audits** separate visual recognition, knowledge application, and evidence integration [14, 15], and controlled hallucinations can escape rubric-based evaluation, motivating **explicit factuality checks** [16].

Supervision quality also depends on **reasoning content and student compatibility**: process supervision provides step-level feedback [17], on-policy distillation studies student-teacher compatibility [9], and ChainPrune [18] reduces redundant reasoning while distinguishing structural efficiency from token length. Building on these observations, MedEvo diagnoses failures with a medical error codebook and admits a repair only if it passes answer, boundary, integrity, and student-relative learnability checks.

### 2.2 Learning from Failures and Localized Repair

Failed trajectories can supply supervision: LEMA trains on mistake-correction pairs [6], TrajFusion combines incorrect and correct trajectories [10], and HHD guides recovery with concise hints [1]. Identifying an error can enable its correction [5], and backtracking distillation and selective regeneration study where correction should begin [19, 20]. MedEvo differs in three respects: the corrected suffix is **generated by the same student** rather than written by a teacher; the auditor **never accesses the reference answer**; and both localization and hints are conditioned on medical tags and error types.



**Figure 3** Overview of our medical multimodal data. Top: distributions across nine modality categories (left) and five difficulty levels (right). Bottom: representative images and functional-test examples illustrating the range of visual inputs. Abbreviations: electrocardiogram (ECG), magnetic resonance cholangiopancreatography (MRCP), positron emission tomography / computed tomography (PET/CT), and magnetic resonance imaging (MRI).


### 2.3 Data Evolution and Recursive Self-Improvement

STaR and ReST<sup>EM</sup> iteratively build supervision from self-generated reasoning filtered by outcome correctness [7, 8], and micro-reflective self-distillation adds capability-dependent trajectory selection [21]. MediSkill-Evo [22] instead evolves validated clinical knowledge banks without updating the backbone. MedEvo follows the self-training loop of STaR and ReST<sup>EM</sup> but differs in two respects: rather than STaR’s rationalization, which supplies the correct answer as a hint [7], it repairs failures with answer-free hints; and it selects each round’s data using **medical-tag and error-type feedback on a fixed feedback set**, linking local recovery to what the student should learn next. Knowledge Navigation couples medically informed local repair with feedback-guided data selection.

See Appendix F for extended related work on medical hallucinations, reasoning supervision, correction methods, and broader self-evolution settings.

## 3 Method

**Knowledge Navigation.** We define **Knowledge Navigation** as a feedback-guided process that uses structured domain knowledge and failure diagnosis to choose **where to intervene, how to recover, and what to learn next**. MedEvo instantiates this process with a medical tag taxonomy and an error codebook: tags organize the clinical context, error types identify the failed reasoning operation, and their associated recovery operations guide intervention. Within a trajectory, the process selects a safe rollback boundary and an answer-free hint while preserving supported reasoning (§3.2). Across rounds, measured capability gaps guide question selection and repair priorities (§3.3). The term navigation therefore describes the link from a structured diagnosis to an actionable decision and its subsequent reassessment. The taxonomy supplies the reference structure; the student’s changing failures determine the direction of learning (Fig. 4).



**Figure 4** MedEvo overview of Knowledge Navigation within trajectories and across training rounds. (a) Medical questions are tagged once by modality, anatomy, disease, and difficulty. (b) The current model generates fresh rollouts for all questions in the first round and feedback-selected questions in later rounds. (c) A gold-blind auditor localizes failures and provides tag-guided minimal hints; the same student preserves a safe prefix and regenerates the suffix. (d) Correct and accepted repaired trajectories update the dataset by question, retaining valid historical targets and replacing them only with accepted improvements, before training the next model from  $M_0$ . (e) Fixed-set evaluation, capability diagnosis, data planning, and targeted question retrieval guide the next rollout round.

*Learning protocol.* A question  $x_i = (I_i, q_i, O_i)$  consists of one or more medical images  $I_i$ , the clinical context and question  $q_i$ , and the answer options  $O_i$ . Following STaR and ReST<sup>EM</sup> [7, 8], each round  $t$  uses the current student  $M_t = \pi_{\theta_t}$  to generate trajectories, collects correct rollouts and accepted repairs into the training set  $D_{t+1}$ , and fine-tunes the next student from the original model  $M_0$  rather than from  $M_t$ . The training pool, feedback set, and test benchmarks are kept separate throughout.

### 3.1 Medical Taxonomy and Failure Diagnosis

#### 3.1.1 Medical Question Tags

We assign each question a four-dimensional hierarchical tag

$$z_i = (z_i^{\text{mod}}, z_i^{\text{anat}}, z_i^{\text{dis}}, z_i^{\text{diff}}), \quad (1)$$

covering modality, anatomy, disease, and difficulty. Tags serve two roles. During repair, the auditor receives only the **observable subset**  $z_i^{\text{obs}}$ , i.e., tags supported by the input itself, excluding answer-derived labels such as the diagnosed disease. Across rounds, tag-level accuracy on the feedback set identifies capability gaps and guides question selection. See Appendix C.1 for the tag hierarchy and annotation rules.



### 3.1.2 Error Types and Their Relation to Repair

We diagnose failures at the process level, building on work that separates medical reasoning from knowledge and audits intermediate steps [14, 15]. Our codebook defines **four parent error types** with five subtypes each, twenty in total (Appendix C.1.2). Each type is paired with a recovery operation used for hint construction (§3.2.2): visual misreading calls for perceptual re-examination; reasoning error for reconstructing the evidence-to-decision link; differential diagnosis error for comparing discriminative features; and knowledge gap for a narrowly scoped knowledge cue. Whereas tags describe the question, error types describe what went wrong in a specific trajectory; because diagnoses vary across rollouts and checkpoints, **each fresh failure is diagnosed independently**.

## 3.2 Localized Repair of Failed Student Trajectories

### 3.2.1 Student Rollouts and Uncertainty-Based Localization

Tyen et al. [5] show that LLMs can correct reasoning errors once the error location is given; we therefore **localize an error before regenerating**. For each selected question,  $M_t$  samples eight trajectories, and answer-correct ones proceed to quality screening (§3.2.4). An incorrect or unparsable trajectory is segmented into semantically complete steps

$$r^- = (s_1, \dots, s_K), \quad (2)$$

such as image observation, feature interpretation, mechanism, differential diagnosis, and option mapping. The goal is to keep the steps before the **first consequential error** and regenerate the rest.

To propose where to audit, we compute the student’s token entropy

$$H_j = - \sum_{v \in \mathcal{V}} \pi_{\theta_t}(v \mid x, r_{<j}^-) \log \pi_{\theta_t}(v \mid x, r_{<j}^-), \quad (3)$$

where  $\mathcal{V}$  is the vocabulary,  $j$  indexes tokens, and  $r_{<j}^-$  is the token prefix. Logs are natural, with  $0 \log 0 = 0$ . For each nonempty step  $s_k$ , let  $\mathcal{J}_k$  be its token-index set and  $H_{s_k} = (H_j)_{j \in \mathcal{J}_k}$ . We combine their **mean and 90th percentile**  $Q_{0.9}$ , and measure the increase over the preceding  $m_k = \min(m, k-1)$  steps:

$$E_k^* = \lambda \frac{1}{|\mathcal{J}_k|} \sum_{j \in \mathcal{J}_k} H_j + (1 - \lambda) Q_{0.9}(H_{s_k}), \quad (4)$$

$$\Delta E_k = \begin{cases} 0, & k = 1, \\ E_k^* - \frac{1}{m_k} \sum_{\ell=k-m_k}^{k-1} E_\ell^*, & 2 \leq k \leq K, \end{cases} \quad (5)$$

For  $k \geq 2$ ,  $m_k > 0$ . After min-max normalizing both profiles (constant profiles map to zero), we select the anchor

$$k^* \in \arg \max_{1 \leq k \leq K} [\tilde{E}_k^* + \tilde{\Delta E}_k], \quad (6)$$

and the audit window covers steps  $k^* \pm 1$ , clipped to valid indices. We use  $\lambda = 0.5$  and  $m = 3$  (Appendix A.2). **Entropy proposes the audit region**; the auditor determines the error location and repair boundary from the images, question, and reasoning context (§3.2.2).

### 3.2.2 Gold-Blind Auditing and Medical Error Diagnosis

Building on hint-guided recovery in HHD [1], an external auditor  $\mathcal{A}$  (GPT-5.6-sol) turns the diagnosis into concise guidance. The auditor is gold-blind: it receives the question and images, the failed trajectory, the entropy proposal  $\mathcal{E}_t(x, r^-)$ , and the observable tags  $z^{\text{obs}}$ , but **no reference answer, reference rationale, or answer-derived metadata**. It returns

$$(a, e, b, h) = \mathcal{A}(x, r^-, \mathcal{E}_t(x, r^-), z^{\text{obs}}), \quad (7)$$

where  $a \in \{\text{REPAIR}, \text{SKIP}\}$  is the action. For REPAIR,  $e$  is the **earliest consequential error** together with its diagnosed type,  $b$  is the last fully valid step before  $e$ , and  $h$  is a **minimal recovery hint**. For SKIP,  $(e, b, h)$  are unused and no continuation is generated. The hint instantiates the recovery operation of the diagnosed type (§3.1.2) in the observable context. For example, a perceptual hint may ask the student to reinspect a region or waveform attribute without stating the finding, and a differential hint may request comparison of a discriminative feature without

naming the diagnosis. Hints must not reveal the answer, an answer option, or a corrected rationale. The audit prompt appears in Appendix D.3.

### 3.2.3 Prefix Preservation and Same-Student Continuation

Like LEMA [6], we train on corrected mistakes, but the correction is generated by the student itself. The boundary splits the failed trajectory into a **preserved prefix and an erroneous suffix**,

$$r^- = p_b \oplus s^-, \quad p_b = (s_1, \dots, s_b). \quad (8)$$

Here  $b \in \{0, \dots, K-1\}$ ,  $p_0$  is empty, and  $\oplus$  denotes token-sequence concatenation. **We keep  $p_b$  verbatim**, append the hint rendered as a natural-language recovery cue  $g(h)$ , and let the same checkpoint continue:

$$s^+ \sim \pi_{\theta_t}(\cdot \mid x, p_b \oplus g(h)), \quad r^* = p_b \oplus g(h) \oplus s^+. \quad (9)$$

The question stays in the user message,  $p_b \oplus g(h)$  prefills an unfinished assistant message, and the structured audit is stored only as metadata. With the prefix fixed, we sample up to five suffixes per failed trajectory, **stopping at the first admitted continuation**, and send answer-correct continuations to admission (§3.2.4; Appendix D.4).

### 3.2.4 Repair Admission and Training Targets

A continuation is admitted only if it passes **answer, boundary, integrity, and learnability checks**. Learnability is scored on the **new suffix alone**, so that a valid prefix cannot mask the difficulty of the correction. Adapting the instruction-following difficulty of Li et al. [23], we measure how much the suffix depends on the task input:

$$D_{\text{repair}} = \frac{\mathcal{L}_t(s^+ \mid x, p_b)}{\mathcal{L}_t(s^+ \mid p_b)}, \quad (10)$$

For a token sequence  $s$  and context  $c$ ,  $\mathcal{L}_t(s \mid c) = -\sum_{j=1}^{|s|} \log \pi_{\theta_t}(s_j \mid c, s_{<j})$ , in nats. Both terms are **computed without the hint**; the denominator omits  $x$  but retains  $p_b$ . A smaller ratio indicates a larger relative loss reduction from explicitly conditioning on  $x$  beyond the prefix, which may itself contain task information. Both suffixes must be nonempty, all losses finite, and the denominator strictly positive. We also measure how strongly the student still prefers its original wrong suffix, using per-token losses to account for length differences:

$$\Delta_{\text{repair}} = \frac{\mathcal{L}_t(s^+ \mid x, p_b)}{|s^+|} - \frac{\mathcal{L}_t(s^- \mid x, p_b)}{|s^-|}. \quad (11)$$

For a repair candidate  $\xi = (x, r^-, r^*, b)$ , the gap is measured in nats per token. A positive value means the student assigns lower average loss to the wrong suffix than to the repair. The learnability gate  $L_\xi$  requires  $D_{\text{repair}} \in [0.35, 0.85]$  and  $\Delta_{\text{repair}} \in [0.02, 0.30]$ , with **thresholds calibrated on the feedback set** and frozen before final evaluation (Appendix A.4). Admission is

$$\text{Accept}(\xi) = C_\xi \wedge A_\xi \wedge N_\xi \wedge L_\xi, \quad (12)$$

where  $C_\xi$  verifies the answer against the reference label,  $A_\xi$  checks the **preserved prefix and boundary**, and  $N_\xi$  rejects unsupported additions and invalid repetition.

Each admitted repair is trained as the complete response  $r^* = p_b \oplus g(h) \oplus s^+$ , including the recovery cue, with **loss on all assistant tokens**:

$$\mathcal{L}_{\text{repair}}(\theta) = -\sum_{j=1}^{|r^*|} \log \pi_{\theta}(r_j^* \mid x, r_{<j}^*). \quad (13)$$

Original correct rollouts use the same objective, and image and user tokens are excluded. If no repair is admitted, we keep a qualified correct rollout or the existing valid target; otherwise no target is added. **No hints or auditor outputs are used at evaluation.**

## 3.3 Feedback-Driven Evolution of Supervision

Across-round Knowledge Navigation follows the iterative self-training pattern of STaR and ReST<sup>EM</sup> [7, 8], but uses **medical-tag and error-type feedback on a fixed set** to allocate generation and repair effort. The cycle is assessment, targeted selection, local repair, screened dataset update, and retraining. The taxonomy and error codebook remain

fixed; **selection priorities and the supervision dataset evolve** as the student changes.

### 3.3.1 First-Round Coverage

The training pool  $\mathcal{T}$  and a separate, fixed feedback set  $V_{\text{feedback}}$  are drawn from public medical datasets (§4.1). In the first round,  $M_0$  generates **rollouts for every question** in  $\mathcal{T}$ ; correct trajectories enter screening, failures are repaired as in §3.2, and accepted targets form  $D_1$ , giving  $M_1 = \text{SFT}(M_0, D_1)$ .

### 3.3.2 Feedback across Medical Tags and Error Types

Before constructing  $D_{t+1}$ , we evaluate  $M_t$  on the feedback set, including  $M_0$  at initialization. For a supported tag group  $g$  with  $n_g > 0$ , the error rate is

$$\epsilon_{g,t} = \frac{1}{n_g} \sum_{i \in V_{\text{feedback},g}} (1 - c_{i,t}), \quad n_g = |V_{\text{feedback},g}|, \quad (14)$$

where  $c_{i,t} \in \{0, 1\}$  indicates correctness. Comparing consecutive checkpoints separates corrections, regressions, persistent errors, and retained successes (Appendix A.5), and diagnosing failed responses yields each group’s error-type distribution. Groups below a support threshold back off to their parent tag. Feedback questions guide selection and threshold calibration but **never enter training**, and test results never influence data construction.

### 3.3.3 Selecting Questions and Updating the Dataset

From the second round on, we allocate the generation budget across groups by the rule-based priority

$$S_{g,t} = \epsilon_{g,t} + \alpha \rho_{g,t}, \quad \rho_{g,t} = \frac{N_{g,t}^{10}}{n_g}, \quad (15)$$

where  $t \geq 1$  and  $N_{g,t}^{10}$  counts correct-to-incorrect transitions from  $M_{t-1}$  to  $M_t$  and  $\alpha \geq 0$  weights regressions. Quotas are proportional to  $S_{g,t}$ , capped by the available unique questions, and rounded deterministically; unused capacity is redistributed, and each multi-tag question is counted once (Appendix A.5). Within each group, candidates are retrieved by tag match and ranked by historical error records. Selected questions receive **fresh rollouts** from  $M_t$ , their failures are diagnosed anew, and failures whose error types match the feedback priorities are repaired first. Feedback thus decides where effort is spent, while the checks in §3.2.4 decide what enters training.

The dataset keeps **at most one active target per question**. An admitted candidate replaces the existing target only if a blinded comparison judges it better in evidence support, reasoning consistency, and conciseness; otherwise the existing target is kept, and all versions are logged separately. Writing each target as a pair  $(x, r)$  keyed by question ID, with  $D_0 = \emptyset$ , replaced targets  $R_t \subseteq D_t$ , accepted replacements  $U_{t+1}$ , and new targets  $A_{t+1}$  for previously unrepresented questions,

$$D_{t+1} = (D_t \setminus R_t) \cup U_{t+1} \cup A_{t+1}, \quad (16)$$

$$M_{t+1} = \text{SFT}(M_0, D_{t+1}). \quad (17)$$

$U_{t+1}$  contains exactly one replacement for each question in  $R_t$ ;  $A_{t+1}$  contains one target per new question, disjoint by question ID from  $D_t$ . Each round restarts from  $M_0$  with a fresh optimizer and learning-rate schedule, so **experience is carried only through the evolving dataset**. Algorithm 1 summarizes the procedure.



**Algorithm 1:** MedEvo Training with Local Repair and Feedback-Guided Data Evolution

**Require:** Initial student  $M_0$ ; training pool  $\mathcal{T}$  with verification labels; fixed feedback set  $V_{\text{feedback}}$ ; gold-blind auditor  $\mathcal{A}$ ; rounds  $R$ ; selection and generation budgets.

$\mathcal{G} \leftarrow \text{AssignMedicalTags}(\mathcal{T}, V_{\text{feedback}})$

$D_0 \leftarrow \emptyset; H \leftarrow \emptyset$

*// Active targets and trajectory history*

**for**  $t \leftarrow 0$  **to**  $R - 1$  **do**

*// Feedback assessment and question selection*

$F_t \leftarrow \text{Evaluate}(M_t, V_{\text{feedback}}, \mathcal{G})$

**if**  $t = 0$  **then**

$Q_t \leftarrow \mathcal{T}$

*// Broad first-round coverage*

**else**

$P_t \leftarrow \text{PlanQuotas}(F_t, F_{t-1}, \mathcal{G})$

$Q_t \leftarrow \text{RetrieveUnique}(\mathcal{T}, P_t, H)$

$D_{t+1} \leftarrow D_t$

*// Fresh student rollouts and local trajectory repair*

**foreach**  $x \in Q_t$  **do**

$\mathcal{R}_x \leftarrow \text{Rollout}(M_t, x, 8)$

$C_x \leftarrow \text{ScreenCorrect}(x, \mathcal{R}_x)$

**foreach** failed  $r^- \in \text{PrioritizeFailures}(\mathcal{R}_x, F_t)$  within budget **do**

$u \leftarrow \text{StepUncertainty}(M_t, x, r^-)$

$(a, e, b, h) \leftarrow \mathcal{A}(x, r^-, u, z^{\text{obs}}(x))$

*// Observable tags only*

$H \leftarrow \text{RecordHistory}(H, x, r^-, a, e, b, h)$

**if**  $a = \text{SKIP}$  **then**

**continue**

$p_b \leftarrow \text{ValidPrefix}(r^-, b)$

**for**  $j \leftarrow 1$  **to** 5 **do**

$s^+ \sim M_t(\cdot \mid x, p_b \oplus g(h))$

$r^* \leftarrow p_b \oplus g(h) \oplus s^+$

$\xi \leftarrow (x, r^-, r^*, b)$

**if**  $\text{Accept}(\xi)$  (Eq. 12) **then**

$C_x \leftarrow C_x \cup \{r^*\}$

**break**

*// Keep one active target per question; retain valid history on ties*

$D_{t+1} \leftarrow \text{RetainOrReplace}(D_{t+1}, x, C_x)$

*// Full-response SFT from the same initial student*

$M_{t+1} \leftarrow \text{SFT}(M_0, D_{t+1})$

*// Reset optimizer and schedule*

**return**  $\{M_t, D_t\}_{t=1}^R$

## 4 Experiments

### 4.1 Experimental Setup

**Models and data.** Our main backbone is Qwen3.5-9B [24]; Qwen3.5-35B and Qwen3.5-397B test scalability (§5.3). The training pool contains 81,293 questions from the training splits of public medical datasets, spanning nine modality categories and five difficulty levels (Fig. 3), and is deduplicated by question and source IDs, patient/case IDs when available, image hashes, and near-duplicate clusters (Appendix C.1.4). A separate, fixed set of 5,000 modality-stratified questions serves as the feedback set and **never enters training**. Dataset

**Table 1** Round-wise repair statistics. Success rate = successfully repaired / sent for repair. Arrows indicate changes from the preceding round.

| Training round             | R1            | R2             | R3             |
|----------------------------|---------------|----------------|----------------|
| Generator                  | R0            | R1             | R2             |
| <b>Sent for repair</b>     | <b>27,233</b> | <b>25,831↓</b> | <b>23,620↓</b> |
| Share of question bank     | 33.50%        | 31.78%         | 29.06%         |
| Successfully repaired      | 19,567        | 18,258         | 16,797         |
| <b>Repair success rate</b> | <b>71.85%</b> | <b>70.68%↓</b> | <b>71.11%↑</b> |
| Mean repair attempts       | 3.37          | 2.48           | 2.12           |

roles and training configurations are given in Appendices B.1 and C.1.5.

*Evaluation.* We use MedUniEval [25], the Hulu-Med implementation built on Lingshu’s MedEvalKit [11], on PMC-VQA [26], VQA-RAD [27], SLAKE [28], PathVQA [29], OmniMedVQA [30], the multimodal track of MedXpertQA (MedXpertQA-MM) [4], and the medical subset of MMMU (MMMU-Med) [31]. All models share the **answering prompt and parser** (Appendix D.1). We report accuracy and the unweighted seven-benchmark mean.

**Table 2** Accuracy (%) on seven medical benchmarks.  denotes medical-specialized models;  denotes trajectory-repair methods (HHD, TrajFusion, and MedEvo). MedEvo combines medical specialization with trajectory repair. Both MedEvo models are evaluated after Round 3. Avg. is the unweighted seven-benchmark mean; MedXQA is the table abbreviation for MedXpertQA-MM; MMMU-Med denotes the medical subset of MMMU. Blue shading highlights MedEvo. Bold and underlining mark the best and second-best values in each column within each model group; ties share the same rank.

| Model                                                                                                                                                                                       | PMC-VQA      | VQA-RAD      | SLAKE        | PathVQA      | OmniMedVQA   | MedXQA       | MMMU-Med     | Avg.         |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| <i>Proprietary models</i>                                                                                                                                                                   |              |              |              |              |              |              |              |              |
| GPT-5.6-sol                                                                                                                                                                                 | <u>60.43</u> | <u>70.29</u> | <b>81.34</b> | <u>54.84</u> | <u>76.67</u> | <u>77.50</u> | <u>86.24</u> | <u>72.47</u> |
| Gemini-3.8-Flash                                                                                                                                                                            | <b>64.52</b> | <b>74.72</b> | <b>81.34</b> | <b>58.58</b> | <b>80.76</b> | <b>86.30</b> | <b>89.10</b> | <b>76.47</b> |
| <i>Open-source models (&lt; 10B)</i>                                                                                                                                                        |              |              |              |              |              |              |              |              |
|  LLaVA-Med-7B [2]                                                                                          | 30.50        | 53.70        | 48.00        | 38.80        | 44.30        | 20.30        | 29.30        | 37.84        |
|  HuatuoGPT-V-7B [32]                                                                                       | 53.30        | 67.00        | 67.80        | 48.00        | 74.20        | 21.60        | 47.30        | 54.17        |
|  BiMediX2-8B [12]                                                                                          | 43.50        | 49.20        | 57.70        | 37.00        | 63.30        | 21.80        | 39.80        | 44.61        |
| InternVL2.5-8B [33]                                                                                                                                                                         | 51.30        | 59.40        | 69.00        | 42.10        | 81.30        | 21.70        | 53.50        | 54.04        |
| InternVL3-8B [34]                                                                                                                                                                           | 53.80        | 65.40        | 72.80        | 48.60        | 79.10        | 22.40        | 59.20        | 57.33        |
|  LINGSHU-7B [11]                                                                                           | 56.30        | 67.90        | <b>83.10</b> | <b>61.90</b> | 82.90        | 26.70        | 54.00        | 61.83        |
| Qwen3.5-9B (Base) [24]                                                                                                                                                                      | 59.20        | 68.10        | 76.80        | 48.80        | 82.20        | 47.40        | 76.50        | 65.57        |
|  Baseline 1 (HHD) [1]                                                                                    | 58.93        | 67.65        | 71.45        | 50.12        | 82.35        | 48.75        | 76.90        | 65.16        |
|  Baseline 2 (TrajFusion) [10]                                                                            | 59.32        | 69.12        | 77.23        | 49.23        | <u>83.46</u> | 49.56        | 77.20        | 66.45        |
| Baseline 3 (Teacher SFT)                                                                                                                                                                    | <u>61.23</u> | <u>70.15</u> | 71.60        | 53.20        | 81.90        | <u>52.30</u> | <u>78.14</u> | <u>66.93</u> |
|   <b>MedEvo (9B)</b>  | <b>62.68</b> | <b>70.40</b> | <u>78.23</u> | <u>53.34</u> | <b>85.41</b> | <b>54.50</b> | <b>79.28</b> | <b>69.12</b> |
| <i>Open-source models (&gt; 10B)</i>                                                                                                                                                        |              |              |              |              |              |              |              |              |
|  HealthGPT-14B [35]                                                                                      | 56.40        | 65.00        | 66.10        | <u>56.70</u> | 75.20        | 24.70        | 49.60        | 56.24        |
|  HuatuoGPT-V-34B [32]                                                                                    | 56.60        | 61.40        | 69.50        | 44.40        | 74.00        | 22.10        | 51.80        | 54.26        |
| InternVL3-14B [34]                                                                                                                                                                          | 54.10        | 66.30        | 72.80        | 48.00        | 78.90        | 23.10        | 63.10        | 58.04        |
| InternVL2.5-38B [33]                                                                                                                                                                        | 57.20        | 61.40        | 70.30        | 46.90        | 79.90        | 24.40        | 61.60        | 57.39        |
| InternVL3-38B [34]                                                                                                                                                                          | 56.60        | 65.40        | 72.70        | 51.00        | 79.80        | 25.20        | 65.20        | 59.41        |
|  LINGSHU-32B [11]                                                                                        | 57.90        | <b>76.50</b> | <b>89.20</b> | <b>65.90</b> | <u>83.40</u> | 30.90        | 62.30        | 66.59        |
| Qwen3.5-35B (Base) [24]                                                                                                                                                                     | <u>62.40</u> | 70.50        | 70.10        | 49.15        | 83.20        | <u>60.00</u> | <u>80.00</u> | <u>67.91</u> |
|   <b>MedEvo (35B)</b> | <b>64.63</b> | <u>75.83</u> | <u>80.30</u> | 56.01        | <b>83.97</b> | <b>66.50</b> | <b>84.02</b> | <b>73.04</b> |

*Baselines.* On Qwen3.5-9B, we compare with two methods that learn from failed trajectories, HHD [1] and TrajFusion [10], and with **teacher-distillation SFT** on rationales generated by Qwen-3.8-Max. External general-purpose and medical models provide additional context (Appendix G).

## 4.2 Main Results

Table 2 reports the results. On Qwen3.5-9B, MedEvo raises mean accuracy from 65.57% to 69.12% and **improves all seven benchmarks**, with the largest gain on MedXpertQA-MM [4] (47.40%  $\rightarrow$  54.50%). It outperforms HHD [1], TrajFusion [10], and teacher-distillation SFT, whose mean accuracies are 65.16%, 66.45%, and 66.93%, respectively. On Qwen3.5-35B, mean accuracy rises from 67.91% to 73.04% (§ 5.3). A question-level paired bootstrap (10,000 resamples; Table 3, right) yields **95% CIs that exclude zero** on every benchmark, including [2.4, 4.7] percentage points (pp) for the mean gain and [5.2, 9.0] for MedXpertQA-MM. These intervals capture question-level uncertainty; variation across training runs is reported in Table 8.

**Table 3** Round-wise accuracy and paired bootstrap gains for the model based on Qwen3.5-9B. Left: accuracy (%) across R0–R3; each round retrains from Base (R0). Bold and underlining mark the best and second-best values per benchmark; arrows compare consecutive rounds, with MedXQA changes in pp. Right: R3 minus R0 and its 95% confidence interval (CI), in pp, from 10,000 question-level paired bootstrap resamples. Avg. is the unweighted seven-benchmark mean; its CI is computed by resampling independently within each benchmark before averaging. All gain intervals exclude zero.

| Benchmark     | Round-wise accuracy |              |                |                     | Paired bootstrap |              |
|---------------|---------------------|--------------|----------------|---------------------|------------------|--------------|
|               | R0                  | R1           | R2             | R3                  | $\Delta$ (R3–R0) | 95% CI       |
| PMC-VQA       | 59.20               | 59.54 ↑      | 60.97 ↑        | <b>62.68</b> ↑      | +3.48            | [+1.9, +5.1] |
| VQA-RAD       | 68.10               | 68.47 ↑      | 68.06 ↓        | <b>70.40</b> ↑      | +2.30            | [+0.8, +3.8] |
| SLAKE         | 76.80               | 76.92 ↑      | 77.34 ↑        | <b>78.23</b> ↑      | +1.43            | [+0.2, +2.7] |
| PathVQA       | 48.80               | 52.38 ↑      | 51.27 ↓        | <b>53.34</b> ↑      | +4.54            | [+2.6, +6.5] |
| OmniMedVQA    | 82.20               | 80.81 ↓      | 81.58 ↑        | <b>85.41</b> ↑      | +3.21            | [+1.5, +4.9] |
| <b>MedXQA</b> | 47.40               | 52.56 ↑ 5.16 | 52.79 ↑ 0.23   | <b>54.50</b> ↑ 1.71 | +7.10            | [+5.2, +9.0] |
| MMMU-Med      | 76.50               | 79.83 ↑      | <b>79.91</b> ↑ | 79.28 ↓             | +2.78            | [+1.3, +4.3] |
| <b>Avg.</b>   | 65.57               | 67.22 ↑      | 67.42 ↑        | <b>69.12</b> ↑      | +3.55            | [+2.4, +4.7] |

### 4.3 Round-wise Evolution

Mean accuracy **increases monotonically across rounds**, from 65.57% to 67.22%, 67.42%, and 69.12% (Table 3, left); §5.2 isolates the contribution of feedback-guided rounds to this improvement. Individual benchmarks fluctuate, e.g., PathVQA [29] (52.38%, 51.27%, 53.34%). Meanwhile, the repair workload falls from 27,233 to 23,620 questions while the success rate stays near 71% (Table 1; Appendix B.2). On the fixed feedback set, every modality group improves from R0 to R3, with accuracy gains ranging from 4.50 to 14.07 pp for functional tests and ophthalmology, respectively; per-round modality and error-type statistics are given in Appendix B.3. Feedback-set trends guide subsequent data selection; the **seven held-out benchmarks measure generalization**.

## 5 Ablation Studies and Analysis

We analyze repair design, feedback across rounds, model scale, and the **learning value of repair supervision**. All analyses except those on scale use Qwen3.5-9B [24].

### 5.1 Repair Design in the First Round

Table 4 compares four first-round repair designs that vary how much of the trajectory is regenerated (full rewrite vs. suffix continuation) and how specific the hint is (none, generic, or type-guided). We evaluate repair success and downstream learning using repair success rate and the resulting model’s mean accuracy. **Type-guided suffix continuation performs best** on both metrics, reaching 71.80% repair success and 67.22% mean accuracy. Generic-hint continuation achieves 64.50% repair success and 64.20% mean accuracy, while hint-guided full rewriting achieves 57.80% and 63.20%, respectively. These comparisons indicate that both prefix preservation and error-type conditioning contribute.

**Table 4** First-round repair designs (9B). Both metrics are percentages; Base mean accuracy is 65.57%. Red ↓ marks lower values than the full method; bold marks the best values.

| Method                                               | Repair success | Mean accuracy |
|------------------------------------------------------|----------------|---------------|
| Hint + full rewrite                                  | 57.80 ↓        | 63.20 ↓       |
| No-hint rewrite                                      | 36.70 ↓        | 61.16 ↓       |
| Generic hint + suffix continuation                   | 64.50 ↓        | 64.20 ↓       |
| <b>Type-guided hint + suffix continuation (Ours)</b> | <b>71.80</b>   | <b>67.22</b>  |

## 5.2 Contribution of Feedback across Rounds

Table 5 isolates the contribution of feedback. All variants start from the same Round-1 checkpoint (67.22%) and run two further rounds: repair without feedback, feedback without medical tags, feedback without error types, and full MedEvo. Full MedEvo reaches 69.12% at R3, compared with 66.67%, 67.53%, and 66.21% for the three variants, respectively. Without feedback or without error-type feedback, accuracy falls below the shared starting checkpoint; without medical tags, it improves only marginally (67.22% to 67.53%). **Combining medical-tag and error-type feedback yields the largest gain in this comparison.**

**Table 5** Feedback ablation (9B). All runs share R1 = 67.22%. Arrows compare mean accuracy with R1;  $\Delta$  is R3 minus R1 in pp.

| Method                      | R2            | R3            | $\Delta$     |
|-----------------------------|---------------|---------------|--------------|
| No feedback (repair only)   | 66.39↓        | 66.67↓        | -0.55        |
| Without medical tags        | 67.32↑        | 67.53↑        | +0.31        |
| Without error-type feedback | 67.01↓        | 66.21↓        | -1.01        |
| <b>Full MedEvo</b>          | <b>67.42↑</b> | <b>69.12↑</b> | <b>+1.90</b> |

## 5.3 Generalization across Model Scales

We apply MedEvo to Qwen3.5-35B for three rounds and to Qwen3.5-397B for two rounds, evaluating on the same seven benchmarks [4, 26, 27, 28, 29, 30, 31] (Table 6). On Qwen3.5-35B, mean accuracy rises from 67.91% to 69.73%, 70.81%, and 73.04%. MedXpertQA-MM [4] peaks at 66.73% in R2 and declines slightly to 66.50% in R3, indicating benchmark-dependent optima. On Qwen3.5-397B, mean accuracy rises from 71.83% to 74.71% and 74.95%, and MedXpertQA-MM from 67.65% to 75.00%. The R2 model improves six of seven benchmarks, with OmniMedVQA [30] essentially unchanged (85.10%  $\rightarrow$  85.05%). At both scales, **mean accuracy increases in every round.**

**Table 6** Round-wise accuracy (%) for models based on Qwen3.5-35B (left) and Qwen3.5-397B (right). R0 is Base; each later round retrains from the same Base. Bold and underlining mark the best and second-best results within each model scale per benchmark. Arrows show changes from the preceding round; numbers beside the highlighted MedXQA results give changes in pp. Avg. is the seven-benchmark mean.

| Benchmark     | 35B   |                |                      |                      | 397B         |                      |                      |
|---------------|-------|----------------|----------------------|----------------------|--------------|----------------------|----------------------|
|               | R0    | R1             | R2                   | R3                   | R0           | R1                   | R2                   |
| PMC-VQA       | 62.40 | <u>63.12</u> ↑ | 60.33↓               | <b>64.63</b> ↑       | 63.47        | <u>64.87</u> ↑       | <b>65.72</b> ↑       |
| VQA-RAD       | 70.50 | 62.08↓         | <u>70.91</u> ↑       | <b>75.83</b> ↑       | 74.72        | <u>76.05</u> ↑       | <b>77.15</b> ↑       |
| SLAKE         | 70.10 | 78.14↑         | <u>79.18</u> ↑       | <b>80.30</b> ↑       | 77.38        | <u>78.89</u> ↑       | <b>79.10</b> ↑       |
| PathVQA       | 49.15 | <u>52.71</u> ↑ | 50.87↓               | <b>56.01</b> ↑       | 52.49        | <b>58.51</b> ↑       | <u>57.90</u> ↓       |
| OmniMedVQA    | 83.20 | <b>85.39</b> ↑ | <u>85.22</u> ↓       | 83.97↓               | <u>85.10</u> | <b>85.15</b> ↑       | 85.05↓               |
| <b>MedXQA</b> | 60.00 | 65.68↑ +5.68   | <b>66.73</b> ↑ +1.05 | <u>66.50</u> ↓ -0.23 | 67.65        | <u>74.40</u> ↑ +6.75 | <b>75.00</b> ↑ +0.60 |
| MMM-U-Med     | 80.00 | 81.02↑         | <u>82.41</u> ↑       | <b>84.02</b> ↑       | 82.03        | <b>85.10</b> ↑       | <u>84.75</u> ↓       |
| <b>Avg.</b>   | 67.91 | 69.73↑         | <u>70.81</u> ↑       | <b>73.04</b> ↑       | 71.83        | <u>74.71</u> ↑       | <b>74.95</b> ↑       |


## 5.4 Learning Value of Repair Supervision

Beyond repair success, we examine three questions about **what makes repair supervision useful** (Fig. 5).


*Which trajectories are worth repairing?* We compare the 8,000 questions with the longest failed trajectories against 8,000 random questions, each used for one round of continuation training. The longest set has lower repair success (55.70% vs. 67.50%) but yields higher downstream accuracy (66.78% vs. 64.35%; Fig. 5a), so trajectories should be selected for **downstream learning value**.

*How many hints help?* All main experiments provide one hint per failed trajectory. Using the longest-8k subset from the previous comparison, we vary the number of hints per failed trajectory from 1 to 3, 5, and 8, keeping five continuation attempts. Mean accuracy rises from 65.78% with one hint to 67.01% with three, then falls to 66.74% and 64.10% with five and eight (Fig. 5b), indicating that **moderate guidance helps** while excessive hints degrade the resulting supervision.

*Which feedback guides evolution?* Prioritizing repairs by training-set errors raises training accuracy from 63.23% to 78.00% but lowers benchmark accuracy from 65.57% to 64.01%, whereas the fixed feedback set yields 71.00%



**Figure 5** Trajectory selection, hint count, and feedback source for the 9B model. (a) Grouped bars compare longest-8k and random-8k selection in single-round training: longer trajectories have lower repair success but higher downstream accuracy. (b) Mean accuracy with 1, 3, 5, or 8 hints per failed trajectory on the longest-8k subset, keeping five continuation attempts; three hints perform best. (c,d) Training accuracy and independent seven-benchmark mean under training-error feedback versus a fixed 5,000-question feedback set. The feedback set guides selection but is excluded from training.



**Figure 6** Analysis across scale, repair budget, entropy, and difficulty. (a) Seven-benchmark mean accuracies for the MedEvo models based on Qwen3.5-9B and Qwen3.5-35B after Round 3, and Qwen3.5-397B after Round 2. (b) Repair acceptance with and without entropy localization; intervals are pointwise 95% Wilson intervals using 45,000 attempts per setting. (c) Erroneous-step rates and step counts by entropy bin; the diagonal provides a numerical reference. (d) Accuracy by medical difficulty, with sample counts on the right axis.

training accuracy and 69.12% benchmark accuracy (Fig. 5c,d). Training failures supply repair candidates, but **an independent feedback set better identifies weaknesses that generalize**.

*Further analyses.* Fig. 6 summarizes additional repair analyses. Entropy localization increases final acceptance, defined as the fraction of repair candidates retained, at every continuation budget: from 19.8–35.8% to 25.2–40.9% for one to eight attempts (Fig. 6b). The erroneous-step rate rises from 16.2% to 64.1% across increasing entropy bins, supporting **entropy as an audit prior** (Fig. 6c). Gains over Base grow with difficulty, from 3.8 to 8.7 pp between L1 and L5 (Fig. 6d).



## 6 Discussion and Insights

We first examine the **reliability of repaired supervision** through a human audit and the **stability of MedEvo** across independent training runs. We then distill what these results, together with the analyses in §5, imply for medical self-evolution, and close with limitations and future directions.

### 6.1 Human Audit

Table 7 reports results for the audited sample of repaired trajectories. Annotators identify 95.5% of submitted trajectories as erroneous ( $\kappa = 0.90$ ) and validate 84.5% of the selected repair boundaries ( $\kappa = 0.85$ ). Hints are answer-blind in 97.2% of cases ( $\kappa = 0.94$ ). Repaired answers are judged correct in 88.3% of cases ( $\kappa = 0.89$ ), and 78.5% of repairs are judged **image-grounded**. Across the four

**Table 7** Human audit pass rates (%) by error family: visual (V), reasoning (R), differential diagnosis (D), and knowledge (K). Agreement:  $\kappa$  with 95% CIs. Blue highlights repaired-output checks.

| Audit criterion                   | Pass (%)    | $\kappa$ | Agreement<br>95% CI | Error family |     |    |    |
|-----------------------------------|-------------|----------|---------------------|--------------|-----|----|----|
|                                   |             |          |                     | V            | R   | D  | K  |
| Original trajectory erroneous     | 95.5        | 0.90     | [0.83, 0.97]        | 96           | 96  | 94 | 96 |
| Valid repair boundary             | 84.5        | 0.85     | [0.66, 0.88]        | 88           | 84  | 82 | 84 |
| No answer leakage in hint         | 97.2        | 0.94     | [0.87, 1.00]        | 96           | 100 | 96 | 97 |
| Image-grounded repaired reasoning | <b>78.5</b> | 0.70     | [0.58, 0.82]        | 76           | 80  | 76 | 82 |
| Correct repaired answer           | <b>88.3</b> | 0.89     | [0.81, 0.97]        | 88           | 89  | 88 | 88 |

error families, answer-correctness pass rates range from 88% to 89%, and grounding pass rates from 76% to 82%.

The audit assesses **answer correctness and visual grounding as complementary dimensions** of repair quality, reporting separate marginal pass rates of 88.3% and 78.5%, respectively. Agreement is higher for answer correctness ( $\kappa = 0.89$ ) than for grounding ( $\kappa = 0.70$ ). These results motivate evaluating both the final answer and the image evidence supporting intermediate deductions.

### 6.2 Stability across Random Seeds

Table 8 reports three independent training runs from the same Base checkpoint. The **improvement is highly reproducible**. Every run follows the ordering Base < R1 < R2 < R3, and R3 averages 69.09%  $\pm$  0.24, compared with 65.57% for Base. Even the weakest run reaches 68.84%, so the improvement far exceeds the variation across seeds. The round-wise pattern is equally consistent: R1→R2 yields a small accuracy gain (0.16–0.22 pp), followed by a larger gain at R2→R3 (1.63–1.71).

Each round retrains from  $M_0$ , so **the evolving dataset carries information across rounds**. Across all three seeds, the R2-to-R3 gain exceeds the R1-to-R2 gain. Longer training runs can characterize the timing and saturation of these gains.

**Table 8** Seven-benchmark mean accuracy (%) over three Qwen3.5-9B runs. SD is the sample standard deviation;  $\Delta$  is gain over Base in pp. Bold/underlining mark best/second-best rounds; R3 is blue.

| Model     | Seed 1       | Seed 2       | Seed 3       | Mean $\pm$ SD                      | $\Delta$               |
|-----------|--------------|--------------|--------------|------------------------------------|------------------------|
| Base      | 65.57        | 65.57        | 65.57        | 65.57 $\pm$ 0.00                   | —                      |
| R1        | 67.22        | 67.05        | 67.38        | 67.22 $\pm$ 0.17                   | $\uparrow$ 1.65        |
| R2        | <u>67.42</u> | <u>67.21</u> | <u>67.60</u> | 67.41 $\pm$ 0.20                   | $\uparrow$ 1.84        |
| <b>R3</b> | <b>69.12</b> | <b>68.84</b> | <b>69.31</b> | <b>69.09 <math>\pm</math> 0.24</b> | $\uparrow$ <b>3.52</b> |

### 6.3 Insights for Medical Self-Evolution

Taken together, our results suggest that effective self-evolution in medical reasoning rests on three choices: **what reasoning to preserve, which repairs to admit** as supervision, and **which failures to revisit next**.

*Defining characteristics of Knowledge Navigation.* Three characteristics connect **medical context, failure diagnosis, and learning decisions**. First, navigation is **knowledge-structured**: tags and error types make a failure interpretable in terms of clinical context and a recoverable reasoning operation. Second, it is **selective and student-relative**: local repair preserves supported steps, and admission assesses whether a correction is suitable supervision for the current student. Third, it is **feedback-guided**: independent assessment redirects subsequent data

selection as weaknesses change. Together, these properties connect the correction of one trajectory to the development of a broader capability. The repair and feedback ablations (Tables 4 and 5) support these design choices within MedEvo.

*Evolution through supervision.* Knowledge Navigation couples two scales of change. **Local recovery** turns an individual failure into a screened training target; **across-round adaptation** uses capability feedback to decide which targets should be generated next. The resulting cycle is assessment, targeted selection, local repair, dataset update, and reassessment after training. In the present implementation, the taxonomy and error codebook stay fixed, while selection priorities, accepted supervision, and the resulting student capabilities change. Each round retrains from  $M_0$  on the updated dataset, so the evolving supervision carries information across rounds. This makes the evolution mechanism explicit and allows regressions and persistent weaknesses to redirect later effort.

*Failed trajectories retain valuable reasoning.* Suffix continuation outperforms full rewriting in both repair success and downstream accuracy (Table 4), supporting the value of **preserving the valid prefix**. In the trajectory-selection experiment, the longest-8k subset has lower repair success than random selection (55.70% vs. 67.50%) and yields higher downstream accuracy (66.78% vs. 64.35%; Fig. 5a). These comparisons support selecting repair candidates by their contribution to subsequent learning.

*Repairability and learning value are different objectives.* Repair success measures whether a correction is reached, whereas downstream accuracy measures what the student learns from it. Our experiments separate the two: harder-to-repair trajectories teach more (Fig. 5a), and guidance beyond a moderate number of hints lowers downstream accuracy (Fig. 5b). Repair should therefore provide minimal guidance and be **evaluated by its effect on subsequent learning**. The same division of labor applies to localization. Entropy reliably indicates where to inspect, with the erroneous-step rate rising from 16.2% to 64.1% across entropy bins (Fig. 6c), while the evidence-grounded audit determines whether an error is present and where regeneration should begin.

*Continued repair needs independent feedback.* Useful repairs change as the student changes. Training-error feedback improves fit to training data while benchmark accuracy declines, whereas the fixed feedback set improves benchmark accuracy across rounds (Fig. 5c,d). Student failures are therefore best used as repair candidates, while **an independent feedback set identifies which capability gaps to target**; medical tags and error types make these gaps actionable.

*Answer correctness and grounding are complementary.* The human audit confirms high answer correctness and shows that visual grounding is a distinct, graded dimension of repair quality (§6.1). Medical self-evolution therefore benefits from assessing **evidence use alongside accuracy** and hint non-leakage.

## 6.4 Limitations and Future Directions

**The current implementation has a fixed knowledge structure.** Its taxonomy and error codebook constrain which weaknesses can be represented, and a strong proprietary auditor supplies the diagnoses. Evaluating open-source auditors and separating the effects of trajectory length from question difficulty would clarify these dependencies. Our benchmarks use multiple-choice and short-answer formats; clinical deployment requires **evaluation in the intended care setting**.

*From diagnosis to grounded recovery.* Future work will investigate **evidence-grounded recovery** through source-attributed retrieval, factuality checks, and image-region supervision. Evaluation should jointly measure answer leakage, grounding, and downstream learning.

*Long-horizon and adaptive navigation.* A research direction is **knowledge-structure evolution**: using persistent failures to propose expert-validated updates to the taxonomy and recovery operations. Longer runs can measure regression, cost, and saturation on independent evaluation data.

*Potential for transferable learning strategies.* Future studies can test the reuse of repair experience across modalities, specialties, and institutions, and its role in allocating annotation and auditor budgets. Key outcomes include transfer performance, rare-condition coverage, and robustness under distribution shift.

## 7 Conclusion

We introduced MedEvo and defined **Knowledge Navigation** as a feedback-guided process that uses structured domain knowledge and failure diagnosis to choose where to intervene, how to recover, and what to learn next.

A gold-blind auditor localizes the first consequential error and provides a minimal, answer-free hint. The same student preserves its valid prefix and regenerates the erroneous suffix, and feedback on a fixed set directs subsequent repairs toward remaining weaknesses. **With SFT alone**, MedEvo improves Qwen3.5-9B, Qwen3.5-35B, and Qwen3.5-397B across seven medical multimodal benchmarks, and on Qwen3.5-9B these gains are reproducible across independent training runs.

Our analyses establish three findings. First, **failed trajectories contain reasoning worth preserving**. Second, **repair success and learning value are distinct objectives**. Third, **independent feedback is essential for sustained improvement**. Together, these findings identify repair selection and feedback-guided targeting as central to effective medical self-evolution.

**Knowledge Navigation links error diagnosis to learning strategy.** Future work can extend this framework through evidence-grounded repair, longer training horizons, and validated updates to the knowledge structure.

*Declaration of Artificial Intelligence (AI) Assistance.* AI tools were used solely for language editing, manuscript formatting, and figure preparation. They were not involved in method development or experimental design. All AI-assisted content was reviewed and revised by the authors, who take full responsibility for the final manuscript.

*Reproducibility Statement.* The method and experimental procedures are described in the main text and appendices. We will release the implementation, prompts, training configurations, evaluation scripts, and data-processing instructions to support reproducibility, subject to the licenses and access restrictions of the underlying datasets.

## References

- [1] Shengjie Wang, Guanghe Li, Zonghan Yang, and Yang Gao. Hindsight hint distillation: Scaffolded reasoning for SWE agents from cot-free answers. *CoRR*, abs/2605.11556, 2026. doi: 10.48550/arxiv.2605.11556. URL <https://dblp.org/rec/journals/corr/abs-2605-11556>.
- [2] Chunyuan Li, Cliff Wong, Sheng Zhang, Naoto Usuyama, Haotian Liu, Jianwei Yang, Tristan Naumann, Hoifung Poon, and Jianfeng Gao. LLaVA-Med: Training a Large Language-and-Vision Assistant for Biomedicine in One Day. In *Advances in Neural Information Processing Systems*, 2023. URL <https://dblp.org/rec/conf/nips/LiWZULYNPG23>.
- [3] Michael Moor, Qian Huang, Shirley Wu, Michihiro Yasunaga, Yash Dalmia, Jure Leskovec, Cyril Zakka, Eduardo Pontes Reis, and Pranav Rajpurkar. Med-Flamingo: a Multimodal Medical Few-shot Learner. In *Machine Learning for Health (ML4H)*, pages 353–367, 2023. URL <https://dblp.org/rec/conf/ml4h/MoorHWYDLZRR23>.
- [4] Yuxin Zuo, Shang Qu, Yifei Li, Zhang-Ren Chen, Xuekai Zhu, Ermo Hua, Kaiyan Zhang, Ning Ding, and Bowen Zhou. Medxpertqa: Benchmarking expert-level medical reasoning and understanding. In *International Conference on Machine Learning*, volume 267, pages 80961–80990, 2025. URL <https://dblp.org/rec/conf/icml/ZuoQLCZH0025>.
- [5] Gladys Tyen, Hassan Mansoor, Victor Carbune, Peter Chen, and Tony Mak. LLMs cannot find reasoning errors, but can correct them given the error location. In *Findings of the Association for Computational Linguistics*, pages 13894–13908, 2024. URL <https://dblp.org/rec/conf/acl/TyenMCCM24>.
- [6] Shengnan An, Zexiong Ma, Zeqi Lin, Nanning Zheng, Jian-Guang Lou, and Weizhu Chen. Learning From Mistakes Makes LLM Better Reasoner. *arXiv preprint arXiv:2310.20689*, 2023. URL <https://arxiv.org/abs/2310.20689>.
- [7] Eric Zelikman, Yuhuai Wu, Jesse Mu, and Noah D. Goodman. STaR: Self-taught reasoner—bootstrapping reasoning with reasoning. In *Advances in Neural Information Processing Systems*, 2022. URL <https://dblp.org/rec/conf/nips/ZelikmanWMG22>.
- [8] Avi Singh, John D. Co-Reyes, Rishabh Agarwal, Ankesh Anand, Piyush Patil, Xavier Garcia, et al. Beyond human data: Scaling self-training for problem-solving with language models. *arXiv preprint arXiv:2312.06585*, 2023. URL <https://arxiv.org/abs/2312.06585>.
- [9] Yaxuan Li, Yuxin Zuo, Bingxiang He, Jinqian Zhang, Chaojun Xiao, Cheng Qian, Tianyu Yu, Huan ang Gao, Wenkai Yang, Zhiyuan Liu, and Ning Ding. Rethinking on-policy distillation of large language models: Phenomenology, mechanism, and recipe. *CoRR*, abs/2604.13016, 2026. URL <https://dblp.org/rec/journals/corr/abs-2604-13016>.
- [10] Jie Deng, Hanshuang Tong, Jun Li, Shining Liang, Ning Wu, Hongzhi Li, and Yutao Xie. Beyond rejection sampling: Trajectory fusion for scaling mathematical reasoning. In Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens, editors, *Findings of the Association for Computational Linguistics, ACL 2026, San Diego, California, United States, July 2-7, 2026*, pages 7943–7959. Association for Computational Linguistics, 2026. doi: 10.18653/V1/2026.FINDINGS-ACL.390. URL <https://doi.org/10.18653/v1/2026.findings-acl.390>.
- [11] LASA Team, Weiwen Xu, Hou Pong Chan, Long Li, Mahani Aljunied, et al. Lingshu: A Generalist Foundation Model for Unified Multimodal Medical Understanding and Reasoning. *CoRR*, abs/2506.07044, 2025. URL <https://dblp.org/rec/journals/corr/abs-2506-07044>.

- [12] Sahal Shaji Mullappilly, Mohammed Irfan Kurpath, Sara Pieri, Saeed Yahya Alseiari, Shanavas Cholakkal, Khaled Aldahmani, Fahad Khan, Rao Anwer, Salman Khan, Timothy Baldwin, and Hisham Cholakkal. Bimedix2: Bio-medical expert llm for diverse medical modalities, 2024. URL <https://arxiv.org/abs/2412.07769>.
- [13] Andrew Sellergren, Sahar Kazemzadeh, Tiam Jaroensri, Atilla Kiraly, Madeleine Traverse, et al. MedGemma Technical Report. *arXiv preprint arXiv:2507.05201*, 2025. doi: 10.48550/arXiv.2507.05201. URL <https://arxiv.org/abs/2507.05201>.
- [14] Rahul Thapa, Qingyang Wu, Kevin Wu, Harrison Zhang, Angela Zhang, Eric Wu, Haotian Ye, Suhana Bedi, Nevin Aresh, Joseph Boen, Shriya Reddy, Ben Athiwaratkun, Shuaiwen Leon Song, and James Zou. Disentangling reasoning and knowledge in medical large language models. *CoRR*, abs/2505.11462, 2025. doi: 10.48550/arxiv.2505.11462. URL <https://dblp.org/rec/journals/corr/abs-2505-11462>.
- [15] Zhaoyang Jiang, Xuanqi Peng, Fei Teng, Zhizhong Fu, Yunsoo Kim, Jiacong Mi, Zicheng Li, and Honghan Wu. Better accuracies, worse reasoning: A step-level audit of medical chain-of-thought distillation. *CoRR*, abs/2605.28301, 2026. doi: 10.48550/ARXIV.2605.28301. URL <https://doi.org/10.48550/arXiv.2605.28301>.
- [16] Griffin Farrow, Lily Sijia Li, Jack Johnson, Tingyan Wang, Philip Torr, William Bolton, and Fabio J. Fehr. When Rubrics Fail: Hallucinations Reveal Blind Spots in Medical AI Evaluation. *CoRR*, abs/2609.12718, 2026. URL <https://arxiv.org/abs/2609.12718>.
- [17] Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In *International Conference on Learning Representations*, 2024. URL <https://dblp.org/rec/conf/iclr/LightmanKBEBLLS24>.
- [18] Weihang Pan, Zhengxu Yu, Yuxiang Zhang, Wenzhi Li, Zhongming Jin, Binbin Lin, Xiaofei He, and Jieping Ye. ChainPrune: Evaluating and Reducing Redundancy in Long Chain-of-Thought Reasoning. *CoRR*, abs/2608.21860, 2026. URL <https://arxiv.org/abs/2608.21860v1>.
- [19] Bing Wang, Shaotian Yan, Chen Shen, Kaiyuan Liu, Sinan Fan, Ximing Li, Rui Miao, Xiaosong Yuan, Zhanming Shen, and Jieping Ye. Backtracking when it strays: Mitigating dual exposure biases in LLM reasoning distillation. *CoRR*, abs/2605.19433, 2026. doi: 10.48550/ARXIV.2605.19433. URL <https://doi.org/10.48550/arXiv.2605.19433>.
- [20] Sophia Xiao Pu, Yumo Xu, Sailik Sengupta, Millennium Bismay, Ruixue Lian, James Gung, Yi-An Lai, and Arshit Gupta. Selective regenerative decoding: Trajectory-level intervention for inference-time reasoning. *CoRR*, abs/2608.24338, 2026. doi: 10.48550/ARXIV.2608.24338. URL <https://doi.org/10.48550/arXiv.2608.24338>.
- [21] Zhilin Huang, Hang Gao, Ziqiang Dong, Yuan Chen, Yifeng Luo, Chujun Qin, Jingyi Wang, Yang Yang, and Guanjuan Jiang. Learning from your own mistakes: Constructing learnable micro-reflective trajectories for self-distillation. *CoRR*, abs/2606.18844, 2026. doi: 10.48550/ARXIV.2606.18844. URL <https://doi.org/10.48550/arXiv.2606.18844>.
- [22] Ruoyu Wu, Shenfu Xie, Yinqian Sun, Haibo Tong, and Feifei Zhao. MediSkill-Evo: Process-Constrained Self-Evolution for Evidence-Grounded Clinical Interaction. *CoRR*, abs/2608.23397, 2026. URL <https://arxiv.org/abs/2608.23397v2>.
- [23] Ming Li, Yong Zhang, Zhitao Li, Jiuhai Chen, Lichang Chen, Ning Cheng, Jianzong Wang, Tianyi Zhou, and Jing Xiao. From quantity to quality: Boosting llm performance with self-guided data selection for instruction tuning. In *Proceedings of the Conference of the North American Chapter of the Association for Computational Linguistics*, 2024. URL <https://dblp.org/rec/conf/naacl/LiZLCC0W0024>.
- [24] Qwen Team. Qwen3.5: Towards native multimodal agents, February 2026. URL <https://qwen.ai/blog?id=qwen3.5>.
- [25] Songtao Jiang, Yuan Wang, Sibao Song, Tianxiang Hu, Chenyi Zhou, Bin Pu, Yan Zhang, Zhibo Yang, Yang Feng, Joey Tianyi Zhou, Jin Hao, Zijian Chen, Ruijia Wu, Tao Tang, Junhui Lv, Hongxia Xu, Hongwei Wang, Jun Xiao, Bin Feng, Fudong Zhu, Kenli Li, Weidi Xie, Jimeng Sun, Jian Wu, and Zuozhu Liu. Hulu-med: A transparent generalist model towards holistic medical vision-language understanding, 2025. URL <https://arxiv.org/abs/2510.08668>.
- [26] Xiaoman Zhang, Chaoyi Wu, Ziheng Zhao, Weixiong Lin, Ya Zhang, Yanfeng Wang, and Weidi Xie. PMC-VQA: Visual Instruction Tuning for Medical Visual Question Answering. *arXiv preprint arXiv:2305.10415*, 2023. URL <https://arxiv.org/abs/2305.10415>.
- [27] Jason J. Lau, Soumya Gayen, Asma Ben Abacha, and Dina Demner-Fushman. A dataset of clinically generated visual questions and answers about radiology images. *Scientific Data*, 5:180251, 2018. doi: 10.1038/sdata.2018.251.
- [28] Bo Liu, Li-Ming Zhan, Li Xu, Lin Ma, Yan Yang, and Xiao-Ming Wu. SLAKE: A Semantically-Labeled Knowledge-Enhanced Dataset for Medical Visual Question Answering. *CoRR*, abs/2102.09542, 2021. URL <https://dblp.org/rec/journals/corr/abs-2102-09542>.
- [29] Xuehai He, Yichen Zhang, Luntian Mou, Eric Xing, and Pengtao Xie. PathVQA: 30000+ Questions for Medical Visual Question Answering. *CoRR*, abs/2003.10286, 2020. URL <https://dblp.org/rec/journals/corr/abs-2003-10286>.
- [30] Yutao Hu, Tianbin Li, Quanfeng Lu, Wenqi Shao, Junjun He, Yu Qiao, and Ping Luo. OmniMedVQA: A New Large-Scale Comprehensive Evaluation Benchmark for Medical LLM. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 22170–22183, 2024. doi: 10.1109/CVPR52733.2024.02093. URL <https://dblp.org/rec/conf/cvpr/HuLLSHQL24>.



- [31] Xiang Yue, Yuansheng Ni, Tianyu Zheng, Kai Zhang, Ruoqi Liu, Ge Zhang, Samuel Stevens, Dongfu Jiang, Weiming Ren, Yuxuan Sun, Cong Wei, Botao Yu, Ruibin Yuan, Renliang Sun, Ming Yin, Boyuan Zheng, Zhenzhu Yang, Yibo Liu, Wenhao Huang, Huan Sun, Yu Su, and Wenhui Chen. MMMU: A Massive Multi-discipline Multimodal Understanding and Reasoning Benchmark for Expert AGI. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 9556–9567, 2023. doi: 10.1109/CVPR52733.2024.00913. URL <https://dblp.org/rec/conf/cvpr/YueNZ0LZSJRSWYY24>.
- [32] Junying Chen, Ruyi Ouyang, Anningzhe Gao, Shunian Chen, Guiming Hardy Chen, Xidong Wang, Ruifei Zhang, Zhenyang Cai, Ke Ji, Guangjun Yu, Xiang Wan, and Benyou Wang. HuatuoGPT-vision, towards injecting medical visual knowledge into multimodal llms at scale. *CoRR*, abs/2406.19280, 2024. doi: 10.48550/ARXIV.2406.19280. URL <https://doi.org/10.48550/arXiv.2406.19280>.
- [33] Zhe Chen, Weiyun Wang, Yue Cao, Yangzhou Liu, Zhangwei Gao, Erfei Cui, Jinguo Zhu, Shenglong Ye, Hao Tian, Zhaoyang Liu, Lixin Gu, Xuehui Wang, Qingyun Li, Yimin Ren, Zixuan Chen, Jiapeng Luo, Jiahao Wang, Tan Jiang, Bo Wang, Conghui He, Botian Shi, Xingcheng Zhang, Han Lv, Yi Wang, Wenqi Shao, Pei Chu, Zhongying Tu, Tong He, Zhiyong Wu, Huipeng Deng, Jiaye Ge, Kai Chen, Min Dou, Lewei Lu, Xizhou Zhu, Tong Lu, Dahua Lin, Yu Qiao, Jifeng Dai, and Wenhui Wang. Expanding performance boundaries of open-source multimodal models with model, data, and test-time scaling. *CoRR*, abs/2412.05271, 2024. doi: 10.48550/ARXIV.2412.05271. URL <https://doi.org/10.48550/arXiv.2412.05271>.
- [34] Jinguo Zhu, Weiyun Wang, Zhe Chen, Zhaoyang Liu, Shenglong Ye, Lixin Gu, Hao Tian, Yuchen Duan, Weijie Su, Jie Shao, Zhangwei Gao, Erfei Cui, Xuehui Wang, Yue Cao, Yangzhou Liu, Xingguang Wei, Hongjie Zhang, Haomin Wang, Weiye Xu, Hao Li, Jiahao Wang, Nianchen Deng, Songze Li, Yinan He, Tan Jiang, Jiapeng Luo, Yi Wang, Conghui He, Botian Shi, Xingcheng Zhang, Wenqi Shao, Junjun He, Yingting Xiong, Wenwen Qu, Peng Sun, Penglong Jiao, Han Lv, Lijun Wu, Kaipeng Zhang, Huipeng Deng, Jiaye Ge, Kai Chen, Limin Wang, Min Dou, Lewei Lu, Xizhou Zhu, Tong Lu, Dahua Lin, Yu Qiao, Jifeng Dai, and Wenhui Wang. InternV3: Exploring advanced training and test-time recipes for open-source multimodal models. *CoRR*, abs/2504.10479, 2025. doi: 10.48550/ARXIV.2504.10479. URL <https://doi.org/10.48550/arXiv.2504.10479>.
- [35] Tianwei Lin, Wenqiao Zhang, Sijing Li, Yuqian Yuan, Binhe Yu, Haoyuan Li, Wanggui He, Hao Jiang, Mengze Li, Xiaohui Song, Siliang Tang, Jun Xiao, Hui Lin, Yueting Zhuang, and Beng Chin Ooi. Healthgpt: A medical large vision-language model for unifying comprehension and generation via heterogeneous knowledge adaptation. In Aarti Singh, Maryam Fazel, Daniel Hsu, Simon Lacoste-Julien, Felix Berkenkamp, Tegan Maharaj, Kiri Wagstaff, and Jerry Zhu, editors, *Forty-second International Conference on Machine Learning, ICML 2025, Vancouver, BC, Canada, July 13-19, 2025*, volume 267 of *Proceedings of Machine Learning Research*. PMLR / OpenReview.net, 2025. URL <https://proceedings.mlr.press/v267/lin25n.html>.
- [36] Junying Chen, Zhenyang Cai, Ke Ji, Xidong Wang, Wanlong Liu, Rongsheng Wang, Jianye Hou, and Benyou Wang. HuatuoGPT-o1, Towards Medical Complex Reasoning with LLMs. *arXiv preprint arXiv:2412.18925*, 2024. URL <https://arxiv.org/abs/2412.18925>.
- [37] Chi Liu, Derek Li, Yan Shu, Robin Chen, Derek Duan, Teng Fang, and Bryan Dai. Fleming-r1: Toward expert-level medical reasoning via reinforcement learning. *CoRR*, abs/2509.15279, 2025. doi: 10.48550/arxiv.2509.15279. URL <https://dblp.org/rec/journals/corr/abs-2509-15279>.
- [38] Haolin Li, Shuyang Jiang, Ruipeng Zhang, Jiangchao Yao, Ya Zhang, and Yanfeng Wang. Eliciting medical reasoning with knowledge-enhanced data synthesis: A semi-supervised reinforcement learning approach. In Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens, editors, *Findings of the Association for Computational Linguistics, ACL 2026, San Diego, California, United States, July 2-7, 2026*, pages 2847–2870. Association for Computational Linguistics, 2026. doi: 10.18653/V1/2026.FINDINGS-ACL.137. URL <https://doi.org/10.18653/v1/2026.findings-acl.137>.
- [39] Jianwen Chen, Xinyu Yang, Peng Xia, Arian Azarang, Yueh Z. Lee, Gang Li, Hongtu Zhu, Yun Li, Beidi Chen, and Huaxiu Yao. Medverse: Efficient and reliable medical reasoning via dag-structured parallel execution. In Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens, editors, *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2026, San Diego, California, United States, July 2-7, 2026*, pages 15320–15336. Association for Computational Linguistics, 2026. doi: 10.18653/V1/2026.ACL-LONG.699. URL <https://doi.org/10.18653/v1/2026.acl-long.699>.
- [40] Yu Zhao, Hao Guan, Yongcheng Jing, Ying Zhang, and Dacheng Tao. Medcog: Maximizing LLM inference density in medical reasoning via meta-cognitive regulation. *CoRR*, abs/2602.07905, 2026. doi: 10.48550/ARXIV.2602.07905. URL <https://doi.org/10.48550/arXiv.2602.07905>.
- [41] Peng Xia, Jinglu Wang, Yibo Peng, Kaide Zeng, Zihan Dong, Xian Wu, Xiangru Tang, Hongtu Zhu, Yun Li, Linjun Zhang, Shujie Liu, Yan Lu, and Huaxiu Yao. MMedAgent-RL: Optimizing Multi-Agent Collaboration for Multimodal Medical Reasoning. In *International Conference on Learning Representations*, 2026. URL <https://dblp.org/rec/journals/corr/abs-2506-00555>.
- [42] Sicheng Yang, Hangjie Yuan, Wenjun Zhang, Jinwang Wang, Yichen Qian, Weihua Chen, Fan Wang, and Lei Zhu. Clinhallu: A benchmark for diagnosing stage-wise hallucinations in medical MLLM reasoning. *CoRR*, abs/2606.14697, 2026. doi: 10.48550/ARXIV.2606.14697.



- 10.48550/ARXIV.2606.14697. URL <https://doi.org/10.48550/arXiv.2606.14697>.
- [43] Honglin Lin, Zheng Liu, Yun Zhu, Chonghan Qin, Juekai Lin, Xiaoran Shang, Conghui He, Wentao Zhang, and Lijun Wu. Mmfinereason: Closing the multimodal reasoning gap via open data-centric methods. *CoRR*, abs/2601.21821, 2026. doi: 10.48550/ARXIV.2601.21821. URL <https://doi.org/10.48550/arXiv.2601.21821>.
- [44] Abhranil Chandra, Ayush Agrawal, Arian Hosseini, Sebastian Fischmeister, Rishabh Agarwal, Navin Goyal, and Aaron Courville. Shape of thought: When distribution matters more than correctness in reasoning tasks. *arXiv preprint arXiv:2512.22255*, 2025.
- [45] Zhanming Shen, Jintao Tong, Shaotian Yan, Chen Shen, Hao Chen, Wentao Ye, Xiaomeng Hu, Rui Miao, Haobo Wang, Junbo Zhao, Gang Chen, and Jieping Ye. Purified OPSD: on-policy self-distillation without losing how to think. *CoRR*, abs/2607.02234, 2026. doi: 10.48550/arxiv.2607.02234. URL <https://dblp.org/rec/journals/corr/abs-2607-02234>.
- [46] Zhinan Hou et al. What Matters in On-Policy Distillation? A Perspective on Data Efficiency and Data Selection. *arXiv preprint arXiv:2609.05198*, 2026. URL <https://arxiv.org/abs/2609.05198>.
- [47] Zixuan Fu et al. Rethinking On-Policy Distillation of Large Language Models II: One Training Example. *arXiv preprint arXiv:2609.04172*, 2026. URL <https://arxiv.org/abs/2609.04172>.
- [48] Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, Shashank Gupta, Bodhisattwa Prasad Majumder, Katherine Hermann, Sean Welleck, Amir Yazdanbakhsh, and Peter Clark. Self-refine: Iterative refinement with self-feedback. In *Advances in Neural Information Processing Systems*, volume 36, pages 46534–46594, 2023. doi: 10.52202/075280-2019. URL <https://dblp.org/rec/conf/nips/MadaanTGHGW0DPY23>.
- [49] Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 36, 2023. URL <https://dblp.org/rec/conf/nips/ShinnCGNY23>.
- [50] Li-Cheng Lan, Andrew Bai, Minhao Cheng, Cho-Jui Hsieh, and Tianyi Zhou. Exploring Expert Failures Improves LLM Agent Tuning. *arXiv preprint arXiv:2504.13145*, 2025. URL <https://arxiv.org/abs/2504.13145>.
- [51] Dongyi Lv et al. FailForge: Distilling Procedural Competence from Persistent Failures into Code Agents. *arXiv preprint arXiv:2608.08570*, 2026. URL <https://arxiv.org/abs/2608.08570>.
- [52] Rajiv Shailesh Chitale, Rahul Madhavan, Taneesh Gupta, Deepanway Ghosal, and Aravindan Raghuvier. Test-time scaling via error localization. *CoRR*, abs/2607.21453, 2026. doi: 10.48550/ARXIV.2607.21453. URL <https://doi.org/10.48550/arXiv.2607.21453>.
- [53] Aviral Kumar, Vincent Zhuang, Rishabh Agarwal, Yi Su, John D. Co-Reyes, Avi Singh, Kate Baumli, Shariq Iqbal, Colton Bishop, Rebecca Roelofs, Lei M. Zhang, Kay McKinney, Disha Shrivastava, Cosmin Paduraru, George Tucker, Doina Precup, Feryal M. P. Behbahani, and Aleksandra Faust. Training language models to self-correct via reinforcement learning. In *International Conference on Learning Representations*, 2025. URL <https://dblp.org/rec/conf/iclr/KumarZASCSBIBRZ25>.
- [54] Aradhya Dixit, Tianxi Liang, and Jai Telang. Project aletheia: Verifier-guided distillation of backtracking for small language models. *arXiv preprint arXiv:2601.14290*, 2026.
- [55] Dmitry Beresnev, Vladimir Makharev, Roman Khalikov, Ivan V. Oseledets, and Petr Anokhin. Search, fail, recover: A training framework for correction-aware reasoning. *CoRR*, abs/2607.07492, 2026. doi: 10.48550/ARXIV.2607.07492. URL <https://doi.org/10.48550/arXiv.2607.07492>.
- [56] Xinbang Dai, Zheyu Xin, Huikang Hu, Lin Ren, Rihui Jin, Guohui Xiao, Kuicai Dong, Zhaocheng Du, and Yuyang Zhang. Evothink: Evolving thinking in large reasoning models via self-pruning and aha-moment preference optimization. In *Proceedings of the Thirty-Fifth International Joint Conference on Artificial Intelligence and the 29th European Conference on Artificial Intelligence, IJCAI/ECAI 2026, Bremen, Germany, 15-21 August 2026*, pages 5693–5701. ijcai.org, 2026. doi: 10.24963/IJCAI.2026/634. URL <https://doi.org/10.24963/ijcai.2026/634>.
- [57] Ziqi Zhao, Xinyu Ma, Liu Yang, Yujie Feng, Daiting Shi, Jingzhou He, Xin Xin, Zhaochun Ren, and Xiao-Ming Wu. ROSD: reflective on-policy self-distillation for language model reasoning across domains. *CoRR*, abs/2605.28014, 2026. doi: 10.48550/ARXIV.2605.28014. URL <https://doi.org/10.48550/arXiv.2605.28014>.
- [58] Yi Duan, Ying Liu, Zirui Tang, Haodong Chen, Jun Zhou, Yumou Liu, et al. The Last AI Built by Humans: Toward Genuine Recursive Self-Improvement. *arXiv preprint arXiv:2609.11873*, 2026. URL <https://arxiv.org/abs/2609.11873>.
- [59] Bo Yin, Qi Li, and Xinchao Wang. On-Policy Self-Evolution via Failure Trajectories for Agentic Safety Alignment. *arXiv preprint arXiv:2605.11882*, 2026. URL <https://arxiv.org/abs/2605.11882>.
- [60] Zhengyang Yan et al. RedFlow: Redirect Failure into Action-Level Corrections for Flow-matching VLA Policy. *arXiv preprint arXiv:2607.27782*, 2026. URL <https://arxiv.org/abs/2607.27782>.
- [61] Ganlong Zhao et al. FLARE: A Failure-Aware Framework for Autonomous Correction and Recovery in Visual-Language Robotic Manipulation. *arXiv preprint arXiv:2608.26645*, 2026. URL <https://arxiv.org/abs/2608.26645>.

- [62] Xiaoyu Luo, Tao Ren, Wenrui Yu, Xiao Li, Qiongxiu Li, and Johannes Bjerva. Capable yet Parsimonious: Extracting and Characterizing Hidden Chain-of-Thought in Frontier Models. arXiv preprint arXiv:2609.26637, 2026. URL <https://arxiv.org/abs/2609.26637>.
- [63] Zheng Liu, Honglin Lin, Xiaoyang Wang, Xin Gao, Yu Li, Mengzhang Cai, Yun Zhu, Zhanping Zhong, Qizhi Pei, Zhuoshi Pan, Xiaoran Shang, Conghui He, Bin Cui, Wentao Zhang, and Lijun Wu. Chartverse: Scaling chart reasoning via reliable programmatic synthesis from scratch. In Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens, editors, Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), ACL 2026, San Diego, California, United States, July 2-7, 2026, pages 7551–7577. Association for Computational Linguistics, 2026. doi: 10.18653/V1/2026.ACL-LONG.344. URL <https://doi.org/10.18653/v1/2026.acl-long.344>.
- [64] Johnny Greco et al. NeMo Data Designer: An Extensible Framework for Multimodal Synthetic Data Generation. arXiv preprint arXiv:2609.17699, 2026. URL <https://arxiv.org/abs/2609.17699>.
- [65] Fanqing Meng, Lingxiao Du, Qiguang Chen, Ziqi Zhao, Haocheng Lu, Mengkang Hu, and Michael Qizhe Shieh. Rsibench-data: Benchmarking data-centric research for recursive self-improvement. CoRR, abs/2607.25886, 2026. doi: 10.48550/ARXIV.2607.25886. URL <https://doi.org/10.48550/arXiv.2607.25886>.
- [66] Tong Zheng et al. Dream-RSI: Recursive Self-Improvement through Evolving Worlds. arXiv preprint arXiv:2609.14858, 2026. URL <https://arxiv.org/abs/2609.14858>.
- [67] Junqi Liu, Yufan He, Yexiao He, Pengfei Guo, Dong Yang, Andriy Myronenko, Can Zhao, Hanrong Ye, Tianhao Qi, Yuyin Zhou, Daguang Xu, and Yucheng Tang. Bat: Towards self-evolving medical research agent with stage rubrics. CoRR, abs/2608.16211, 2026. doi: 10.48550/ARXIV.2608.16211. URL <https://doi.org/10.48550/arXiv.2608.16211>.

## Appendix Contents and Reading Guide

This guide links to each supplementary section and its starting page.

| Section             | Contents                                                                                                                                                                                                                        | Page               |
|---------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------|
| <a href="#">A</a>   | <b>Method Details</b><br>Algorithms, mathematical formulation, and implementation of navigation and local repair.                                                                                                               | <a href="#">22</a> |
| <a href="#">A.1</a> | <b>Failure Definition and Step Segmentation</b><br>Failed answers, semantic steps, and prefix boundaries.                                                                                                                       | <a href="#">22</a> |
| <a href="#">A.2</a> | <b>Entropy-Based Localization</b><br>Step scores, audit windows, and proposal assessment.                                                                                                                                       | <a href="#">22</a> |
| <a href="#">A.3</a> | <b>Gold-Blind Audit and Recovery Operators</b><br>Auditor inputs, recovery hints, and cue rendering.                                                                                                                            | <a href="#">22</a> |
| <a href="#">A.4</a> | <b>Continuation, Admission, and Training Targets</b><br>Budgets, learnability gates, admission, and fallback.                                                                                                                   | <a href="#">23</a> |
| <a href="#">A.5</a> | <b>Feedback Statistics, Quota Allocation, and Dataset Update</b><br>Transitions, quotas, retrieval, and target replacement.                                                                                                     | <a href="#">23</a> |
| <a href="#">A.6</a> | <b>Settings of Analysis Experiments</b><br>Repair, feedback, hint-count, and entropy comparisons.                                                                                                                               | <a href="#">24</a> |
| <a href="#">B</a>   | <b>Dataset Statistics</b><br>Dataset sizes and roles, round-wise construction counts, and detailed feedback-set measurements.                                                                                                   | <a href="#">24</a> |
| <a href="#">B.3</a> | <b>Additional Experimental Details</b><br>Evolution of twenty error subtypes, plus round-wise modality and error-family statistics.                                                                                             | <a href="#">25</a> |
| <a href="#">B.4</a> | <b>Detailed Error Statistics for the Current Rounds</b><br>Numerical rates for all twenty subtypes and repair attempts, successes, and success rates by error family.                                                           | <a href="#">26</a> |
| <a href="#">C</a>   | <b>Medical Taxonomy and Error Categories</b><br>Medical tags, error definitions, information-access rules, and reproducibility specifications.                                                                                  | <a href="#">27</a> |
| <a href="#">C.1</a> | <b>Taxonomy, Permissions, and Execution Contracts</b><br>Four-axis tag dictionary; four error families and twenty subtypes; access matrix; structured records and plan validation; training configuration and evaluation scope. | <a href="#">27</a> |
| <a href="#">D</a>   | <b>Prompts</b><br>Templates for answering questions, planning feedback-driven rounds, auditing failures, and continuing repaired trajectories.                                                                                  | <a href="#">32</a> |
| <a href="#">D.1</a> | <b>Shared Question-Answering Prompt</b><br>The common question-answering instruction.                                                                                                                                           | <a href="#">32</a> |
| <a href="#">D.2</a> | <b>Main-Agent Feedback and Next-Round Planning</b><br>Inputs, constraints, and structured outputs for diagnosing weaknesses and planning the next training round.                                                               | <a href="#">32</a> |
| <a href="#">D.3</a> | <b>Stage 1 Audit Prompt</b><br>Failure diagnosis, localization, and construction of an error-conditioned recovery cue.                                                                                                          | <a href="#">33</a> |
| <a href="#">D.4</a> | <b>Stage 2 Continuation Prompt</b><br>Same-student continuation from the retained prefix and recovery cue.                                                                                                                      | <a href="#">35</a> |
| <a href="#">E</a>   | <b>Case Study</b><br>Worked radiograph examples showing the question, image, retained prefix, hint, and continuation.                                                                                                           | <a href="#">35</a> |
| <a href="#">E.1</a> | <b>Case 08: Chest radiograph</b><br>Chest-image example of local trajectory repair.                                                                                                                                             | <a href="#">35</a> |
| <a href="#">E.2</a> | <b>Case 31: Elbow radiograph</b><br>Elbow-image example of local trajectory repair.                                                                                                                                             | <a href="#">37</a> |
| <a href="#">F</a>   | <b>Additional Related Work</b><br>Extended medical-AI background, hallucinations, supervision quality, correction methods, and broader self-evolution settings.                                                                 | <a href="#">39</a> |
| <a href="#">G</a>   | <b>Additional Baseline Details</b><br>Backbones, comparison models, trajectory-supervision baselines, and the evaluation protocol.                                                                                              | <a href="#">40</a> |

## A Method Details

This appendix provides implementation details that complement §3. Within round  $t$ , **the generator**  $M_t = \pi_{\theta_t}$  **is frozen**: it produces rollouts, entropy profiles, continuations, and learnability scores, while the next student is trained from  $M_0$  (Eq. 17).

### A.1 Failure Definition and Step Segmentation

Given a question  $x = (I, q, O)$  with reference answer  $y \in O$ ,  $M_t$  samples a trajectory  $r$  and a parsed answer  $\hat{y}$ . A trajectory  $r^-$  fails if  $\hat{y} \neq y$  or no valid answer can be parsed. Each failed trajectory is segmented into  $K$  **semantically complete steps** (Eq. 2), typically covering image observation, feature interpretation, mechanism, differential diagnosis, and option mapping. A boundary  $b \in \{0, \dots, K-1\}$  splits  $r^-$  into the preserved prefix  $p_b = (s_1, \dots, s_b)$  and the erroneous suffix  $s^- = (s_{b+1}, \dots, s_K)$  (Eq. 8), where  $b = 0$  denotes an empty prefix.

### A.2 Entropy-Based Localization

*Step scores.* Token entropy  $H_j$ , the step score  $E_k^*$ , and the local surge  $\Delta E_k$  follow Eqs. 3–5, with entropy measured in nats. We use  $\lambda = 0.5$  and  $m = 3$ , with  $m_k = \min(m, k-1)$  for  $k > 1$  and  $\Delta E_1 = 0$ . Combining the mean with the 90th percentile  $Q_{0.9}$  balances average uncertainty with an upper-tail signal within each step.

*Anchor and window.* Both profiles are **min–max normalized** within each trajectory, giving  $\tilde{E}_k^*$  and  $\tilde{\Delta E}_k$ ; a constant profile maps to zero. The anchor is  $k^* \in \arg \max_{1 \leq k \leq K} (\tilde{E}_k^* + \tilde{\Delta E}_k)$ , and the audit window is

$$\mathcal{W}(k^*) = \{s_{\max(1, k^* - w_\ell)}, \dots, s_{\min(K, k^* + w_r)}\}, \quad w_\ell = w_r = 1. \quad (18)$$

The step preceding the window defines a provisional boundary

$$b_0 = \max(0, k^* - w_\ell - 1). \quad (19)$$

Both  $\mathcal{W}(k^*)$  and  $b_0$  are proposals; **the auditor sets the final boundary**  $b$  (Appendix A.3).

*Proposal assessment.* For every audited trajectory, the auditor classifies the entropy proposal as HIT (the first consequential error lies at or adjacent to  $k^*$ ), PARTIAL\_HIT (the window covers the propagated error, but the unsupported premise begins earlier), MISS (the first error lies outside the window), or NON\_ERROR\_UNCERTAINTY (the high-entropy region is medically and logically valid). These labels are logged with each repair record and support the analysis in Fig. 6c.

### A.3 Gold-Blind Audit and Recovery Operators

*Auditor.* We use GPT-5.6-sol as the auditor  $\mathcal{A}$  in Eq. 7. Its inputs are the images, the question and options, the failed trajectory, the entropy profile ( $E_k^*$ , normalized scores,  $\Delta E_k$ ,  $k^*$ , and  $\mathcal{W}(k^*)$ ), and **observable tags with provenance**. The reference answer, reference rationale, answer-derived tags, and retrieval metadata are **excluded through allowlist serialization** (Appendix C.1.3). The auditor returns a structured record (Appendix D.3) containing an action in  $\{\text{REPAIR}, \text{SKIP}\}$ , the first consequential error  $e$  with its type and subtype, the boundary  $b$  with a verbatim ending span (from  $s_b$  if  $b > 0$ , empty if  $b = 0$ ), and the hint  $h$ . It returns SKIP when only answer formatting or option selection is wrong and no reasoning revision is justified. The final boundary may differ from  $b_0$ . The auditor **traces evidence dependencies backward** and moves the boundary earlier when a later entropy surge stems from an upstream unsupported premise; conversely, it retains a high-entropy step that is medically and logically valid.

*Recovery operators.* An error type selects a recovery operator  $o$ ; the text hint  $h$  instantiates that operator in natural language. The operator belongs to

$$o \in \{\mathcal{P}_{\text{inspect}}, \mathcal{R}_{\text{reconstruct}}, \mathcal{D}_{\text{disambiguate}}, \mathcal{K}_{\text{bridge}}\}, \quad (20)$$

each paired with a parent error type (§3.1.2); subtype-level operations are listed in Appendix C.1.2. Operators specify the cognitive operation required for recovery.

- **Perceptual re-examination** ( $\mathcal{P}_{\text{inspect}}$ ) names a visual region or attribute to revisit, such as lesion distribution, boundary shape, waveform morphology, stain localization, or cross-view consistency, without stating the finding.

- **Inference reconstruction** ( $\mathcal{R}_{\text{reconstruct}}$ ) asks the student to rebuild the chain from evidence to mechanism, diagnosis, and the requested decision, for example when supportive and causal treatments are confused or care steps are misordered.
- **Differential disambiguation** ( $\mathcal{D}_{\text{disambiguate}}$ ) requests comparison along the most discriminative feature, such as contrast pattern, onset profile, lesion location, signature pathology, or an exclusion criterion, without naming the diagnosis.
- **Knowledge bridging** ( $\mathcal{K}_{\text{bridge}}$ ) supplies a narrowly scoped concept, such as a diagnostic criterion, drug mechanism, or immunohistochemical localization rule. It may state the concept but must not map it to an answer option.

*Rendering.* The cue  $g(h)$  consists of a natural transition followed by the hint, with flexible wording. Error labels and audit explanations are **stored as metadata** and are never shown to the student.

#### A.4 Continuation, Admission, and Training Targets

*Continuation budget.* The current checkpoint generates eight rollouts per selected question. Each failed trajectory selected for repair is audited once and receives a single hint. With  $p_b \oplus g(h)$  fixed, up to five suffixes are sampled (Eq. 9); **sampling stops at the first admitted continuation**. The user message contains  $x$ , the unfinished assistant message contains  $p_b \oplus g(h)$ , and generation continues within the same message **without repeating or modifying the prefix** (Appendix D.4).

*Learnability gate.* Using the suffix scores  $D_{\text{repair}}$  and  $\Delta_{\text{repair}}$  (Eqs. 10–11), computed under  $M_t$  **without the hint**,

$$L_\xi = \mathbb{1}[D_{\min} \leq D_{\text{repair}}(\xi) \leq D_{\max} \wedge \Delta_{\min} \leq \Delta_{\text{repair}}(\xi) \leq \Delta_{\max}], \quad (21)$$

with  $[D_{\min}, D_{\max}] = [0.35, 0.85]$  and  $[\Delta_{\min}, \Delta_{\max}] = [0.02, 0.30]$ . The lower bound on  $\Delta_{\text{repair}}$  requires that the student still prefers its original wrong suffix, so the repair carries information the student has not yet acquired. The upper bound limits the increase in per-token loss relative to the wrong suffix. Thresholds are selected on the feedback set and **frozen before final evaluation**. Feedback-set questions are **excluded from training targets**. Thresholds, retained fractions, and generator hashes are logged with each repair record (Appendix C.1.4).

*Admission checks.* Admission follows Eq. 12.  $C_\xi$  compares the parsed answer with the reference label.  $A_\xi$  **verifies that the exported prefix matches** the auditor’s verbatim span and that  $b$  precedes the diagnosed error.  $N_\xi$  **rejects continuations** that introduce symptoms, history, measurements, or findings absent from the input, as well as repetition loops. Candidates with either suffix empty, nonfinite losses, or a nonpositive denominator in Eq. 10 are rejected.

*Training targets and fallback.* Admitted repairs are exported as  $r^* = p_b \oplus g(h) \oplus s^+$  and trained with Eq. 13; original correct rollouts and historical repairs use the same **full-response loss**, and image and user tokens are masked. If no repair is admitted for a question, a quality-screened correct rollout or the existing valid target is retained; if neither exists, the question contributes no target. **No hint or auditor output is used at evaluation**.

#### A.5 Feedback Statistics, Quota Allocation, and Dataset Update

*Feedback statistics.* Besides the error rate in Eq. 14, consecutive checkpoints yield transition counts

$$N_{g,t}^{uv} = \sum_{i \in V_{\text{feedback},g}} \mathbb{1}[c_{i,t-1} = u \wedge c_{i,t} = v], \quad u, v \in \{0, 1\}, \quad t \geq 1, \quad (22)$$

which separate **corrected errors** ( $N^{01}$ ), **regressions** ( $N^{10}$ ), **persistent errors** ( $N^{00}$ ), and **retained successes** ( $N^{11}$ ). The priority in Eq. 15 normalizes regressions by  $n_g$ . For diagnostics, we also report the conditional regression rate  $N_{g,t}^{10}/(N_{g,t}^{10} + N_{g,t}^{11})$  only when its denominator is positive; it is undefined and omitted when the preceding checkpoint has no correct examples. Groups comprise supported single-axis tags and predeclared intersections. A group with  $n_g < n_{\min}$  backs off to its nearest supported ancestor, and identical resolved nodes are merged. Each question is counted once within a group, and overall statistics are computed over the deduplicated union of questions.

*Quota allocation.* From the second round on, the budget  $B_t$  is allocated in proportion to  $S_{g,t}$  (Eq. 15), using largest-remainder rounding with deterministic tie-breaking. Quotas are **capped by the available unique training questions**, and surplus capacity is **redistributed among eligible groups**. Each multi-tag question is charged to the budget once. If all scores are zero, questions are sampled uniformly from eligible unique questions. Token and



auditor-call caps impose additional limits, and any unfilled quota is recorded. The priority and quota rule is Eq. 15; the planning template in Appendix D.2 assembles a plan and tag-based retrieval requests subject to these numerical constraints. The validator in Appendix C.1.4 checks the plan before execution.

*Retrieval and repair priority.* Within each group, candidates are retrieved by tag match and ranked by historical error records. Candidates may include questions already used in training. Selected questions receive **fresh roll-outs** from  $M_t$ , and their failures are diagnosed anew. Failures matching both the modality and the error type prioritized by feedback are repaired first, followed by modality-only or error-type-only matches. Matching uses **current diagnoses**. Priority determines where repair effort is spent; admission (Appendix A.4) determines what enters training.

*Dataset update.* Questions are deduplicated by canonical ID, and each keeps **at most one active target**. An admitted candidate replaces the existing target only if a comparison auditor, blinded to round and model labels, judges it better in evidence support, reasoning consistency, and conciseness; **ties retain the existing target**. All versions and rollout-specific diagnoses are logged in a separate history. Each round retrains from  $M_0$  on the complete updated dataset (Eqs. 16–17) with a fresh optimizer and learning-rate schedule. Generation caps are fixed in advance, and plan records follow Appendix C.1.4.

## A.6 Settings of Analysis Experiments

Unless stated otherwise, analyses use Qwen3.5-9B with the training configuration in Table 14 and the evaluation protocol in §4.1.

*Repair designs (Table 4).* All variants are applied in the first round under the same rollout and continuation budgets. Hint + full rewrite uses the type-guided hint but regenerates the trajectory from the start of the response. No-hint rewrite regenerates the full trajectory without a hint. Generic hint + suffix continuation keeps the auditor’s boundary but replaces the type-guided hint with a fixed generic cue.

*Feedback ablations (Table 5).* All variants start from the shared R1 checkpoint and use the same budgets. No feedback omits feedback-guided question selection. Without medical tags omits medical-tag feedback. Without error-type feedback uses tag-group quotas but does not match repairs to prioritized error types.

*Trajectory selection (Fig. 5a).* The longest-8k subset contains the 8,000 questions whose failed trajectories are longest; the random-8k subset is sampled uniformly. Each subset is used for one round of continuation training from  $M_0$ .

*Hint count (Fig. 5b).* On the longest-8k subset, each failed trajectory receives  $k \in \{1, 3, 5, 8\}$  hints, with five continuation attempts.

*Feedback source (Fig. 5c,d).* Training-error feedback uses errors from the training data to guide selection instead of errors on the feedback set.

*Retry budget and entropy (Fig. 6b).* We vary the continuation budget over  $\{1, 3, 5, 8\}$  with and without entropy localization. Final acceptance is the fraction of repair candidates admitted, computed over 45,000 attempts per setting with 95% Wilson intervals.

*Entropy bins and difficulty (Fig. 6c,d).* The analysis compares erroneous-step rates across entropy bins. Difficulty levels follow the rubric in Appendix C.1.1.

## B Dataset Statistics

### B.1 Dataset Roles and Evaluation Scope

The study separates the training bank, the fixed feedback set, and the final evaluation benchmarks. Their roles differ: training questions supply supervision, feedback questions guide diagnosis and data selection, and benchmark questions measure downstream performance.

The training bank is constructed from public medical training splits and deduplicated using question and source identifiers, available patient or case identifiers, image hashes, and near-duplicate clusters. See Appendix C.1.4 for more details of the stored records and information-access rules. Fig. 3 shows nine modality categories and five difficulty levels; the tag hierarchy is described in Appendix C.1.

**Table 9** Dataset roles and reported sizes. Test coverage is reported as the number of benchmarks.

| Resource           | Size             | Role                                                                                |
|--------------------|------------------|-------------------------------------------------------------------------------------|
| Training bank      | 81,293 questions | Source of original correct rollouts, accepted repairs, and retained valid targets.  |
| Fixed feedback set | 5,000 questions  | Medical capability diagnosis and subsequent data selection; excluded from training. |
| Final evaluation   | 7 benchmarks     | Independent downstream accuracy reporting.                                          |

Evaluation uses MedUniEval from Hulu-Med [25], built on Lingshu’s MedEvalKit [11]. The seven benchmarks are PMC-VQA [26], VQA-RAD [27], SLAKE [28], PathVQA [29], OmniMedVQA [30], MedXpertQA-MM [4], and the medical subset of MMMU [31]. Their accuracies are averaged without weighting by benchmark size. Modality and error-family results below describe the feedback set, separately from benchmark evaluation.

## B.2 Round-wise Data Construction

Table 1 highlights a declining repair workload across training rounds. The number of questions submitted for repair falls from **27,233 to 23,620**, a reduction of **3,613 questions (13.27%)** relative to R1. Their share of the fixed question bank decreases from 33.50% to 29.06%. The unit of this workload measure is a question submitted to the repair stage.

The final training set remains fixed at **81,293 questions** in each round, combining qualified original correct trajectories with **accepted repairs** and retained valid historical targets. Thus, repair workload decreases while training coverage stays constant. The number of successfully repaired questions decreases from 19,567 to 16,797 as fewer questions are submitted. Repair success rates are 71.85%, 70.68%, and 71.11%, remaining stable across rounds. With at most five continuation attempts per repair, the mean number of attempts also falls from **3.37 to 2.12**, indicating improved hint-guided correction efficiency. Each column identifies the generator checkpoint and the subsequent training round.

## B.3 Additional Experimental Details

### B.3.1 Evolution across Twenty Error Subtypes

Fig. 7 expands the four families into twenty subtypes. Bar heights are normalized by each subtype’s R0 rate, while in-bar labels retain the reported percentages. All subtypes are below R0 at R3. The knowledge-gap family has an R3/R0 ratio of approximately 0.89, compared with 0.72–0.75 for the other families, indicating a smaller relative reduction in the reported breakdown, conditional on a consistent counting denominator across rounds. Rates are reported separately by error subtype.

### B.3.2 Round-wise Feedback Statistics


Tables 10 and 11 report the per-round modality error rates and error-family percentages; §4.3 summarizes the modality trends. Blue downward arrows indicate decreases, red upward arrows increases, and equality signs unchanged rates. The error-family percentages are reported separately from the modality-level error rates.

**Table 10** Modality error rates (%) on the fixed feedback set; N is group size. Blue shading marks R3. Arrows compare consecutive rounds; bold/underlining mark the two lowest rates.

| Modality         | N    | R0    | R1           | R2             | R3               |
|------------------|------|-------|--------------|----------------|------------------|
| Radiology        | 1509 | 37.18 | 35.98        | ↓ <b>31.28</b> | ↓ <u>32.21</u> ↑ |
| Nuclear medicine | 234  | 34.62 | 26.50        | ↓ <u>24.79</u> | ↓ <b>22.65</b> ↓ |
| Ultrasound       | 351  | 44.73 | <u>39.03</u> | ↓ 39.60 ↑      | ↑ <b>36.18</b> ↓ |
| Endoscopy        | 247  | 32.39 | 30.36        | ↓ <u>27.13</u> | ↓ <b>21.86</b> ↓ |
| Pathology        | 378  | 34.13 | 34.66        | ↑ 30.95        | ↓ <b>28.31</b> ↓ |
| Ophthalmology    | 135  | 34.07 | 27.41        | ↓ <u>22.96</u> | ↓ <b>20.00</b> ↓ |
| Dermatology      | 593  | 30.35 | 27.82        | ↓ <u>24.62</u> | ↓ <b>23.10</b> ↓ |
| Dental imaging   | 308  | 40.26 | <u>37.01</u> | ↓ 37.66 ↑      | ↑ <b>32.79</b> ↓ |
| Functional tests | 1245 | 20.00 | 17.75        | ↓ <u>16.22</u> | ↓ <b>15.50</b> ↓ |

**Table 11** Error-family percentages, summed over five subtypes per family. These values are not directly comparable to the modality error rates. Blue shading marks R3; arrows and bold/underlining follow Table 10.

| Error family       | R0    | R1    | R2             | R3               |
|--------------------|-------|-------|----------------|------------------|
| Visual misreading  | 14.10 | 13.79 | ↓ <u>11.40</u> | ↓ <b>10.23</b> ↓ |
| Reasoning error    | 12.50 | 12.20 | ↓ <u>10.40</u> | ↓ <b>9.40</b> ↓  |
| Differential error | 5.40  | 5.25  | ↓ <u>4.20</u>  | ↓ <b>3.90</b> ↓  |
| Knowledge gap      | 2.19  | 2.49  | ↑ <u>2.16</u>  | ↓ <b>1.96</b> ↓  |



**Figure 7** Twenty error subtypes of the 9B model across R0–R3. Each panel contains five subtypes. Bar heights divide reported rates by the corresponding subtype’s R0 rate; in-bar labels give original percentages. Light-blue arrows and multipliers show R3 relative to R0, with lower values indicating larger relative reductions. Family R3/R0 is the ratio of the summed subtype rates. Horizontal codes such as R1–R5 denote reasoning-error subtypes, distinct from round labels R0–R3 in the legend.

## B.4 Detailed Error Statistics for the Current Rounds

**Table 12** Error-subtype rates (%): visual/reasoning subtypes (left) and differential/knowledge subtypes (right). Blue shading marks R3. Subtype codes R1–R5 denote reasoning errors, not training rounds. Arrows compare each rate with the preceding round: blue ↓ marks a decrease, red ↑ an increase, and = no change. Bold and underlining mark the lowest and second-lowest distinct rates in each row.

| Subtype                  | R0   | R1     | R2            | R3            | Subtype                    | R0          | R1            | R2            | R3            |
|--------------------------|------|--------|---------------|---------------|----------------------------|-------------|---------------|---------------|---------------|
| V1 Missed key finding    | 3.12 | 2.95 ↓ | <u>2.36</u> ↓ | <b>2.12</b> ↓ | D1 Missing candidate       | 1.21        | 1.05 ↓        | <u>0.92</u> ↓ | <b>0.86</b> ↓ |
| V2 Spurious finding      | 2.87 | 2.82 ↓ | <u>2.28</u> ↓ | <b>2.05</b> ↓ | D2 No discriminating cue   | 1.08        | 1.12 ↑        | <u>0.88</u> ↓ | <b>0.82</b> ↓ |
| V3 Location / laterality | 2.65 | 2.68 ↑ | <u>2.20</u> ↓ | <b>1.98</b> ↓ | D3 Premature closure       | 1.14        | 0.98 ↓        | <u>0.84</u> ↓ | <b>0.78</b> ↓ |
| V4 Sequence / phase      | 2.91 | 2.74 ↓ | <u>2.32</u> ↓ | <b>2.08</b> ↓ | D4 Candidate ranking       | 0.97        | 1.08 ↑        | <u>0.80</u> ↓ | <b>0.74</b> ↓ |
| V5 Cross-image alignment | 2.55 | 2.60 ↑ | <u>2.24</u> ↓ | <b>2.00</b> ↓ | D5 Exclusion misuse        | 1.00        | 1.02 ↑        | <u>0.76</u> ↓ | <b>0.70</b> ↓ |
| R1 Omitted evidence      | 2.88 | 2.55 ↓ | <u>2.16</u> ↓ | <b>1.95</b> ↓ | K1 Sign interpretation     | 0.51        | <u>0.48</u> ↓ | <u>0.48</u> = | <b>0.44</b> ↓ |
| R2 Conflicting evidence  | 2.41 | 2.42 ↑ | <u>2.08</u> ↓ | <b>1.88</b> ↓ | K2 Diagnostic criteria     | <u>0.44</u> | 0.52 ↑        | <u>0.44</u> ↓ | <b>0.40</b> ↓ |
| R3 Unsupported inference | 2.67 | 2.48 ↓ | <u>2.12</u> ↓ | <b>1.92</b> ↓ | K3 Anatomy / mechanism     | 0.42        | 0.45 ↑        | <u>0.40</u> ↓ | <b>0.36</b> ↓ |
| R4 Temporal / numeric    | 2.22 | 2.35 ↑ | <u>2.04</u> ↓ | <b>1.85</b> ↓ | K4 Conditions / exceptions | 0.40        | 0.50 ↑        | 0.42 ↓        | <b>0.38</b> ↓ |
| R5 Answer mapping        | 2.32 | 2.40 ↑ | <u>2.00</u> ↓ | <b>1.80</b> ↓ | K5 Test / treatment rule   | <u>0.42</u> | 0.54 ↑        | <u>0.42</u> ↓ | <b>0.38</b> ↓ |

Table 13 separates repair volume from success rate. Counts measure attempted and successful repairs within each error family.

Reasoning-error repair success rises from 56.43% at R1 to 83.18% at R3, while knowledge-gap repair success falls from 90.28% to 48.82%. These rates describe the repair candidates submitted in each round.

**Table 13** Repair counts and success rates by error family. Gray arrows track counts; blue/red arrows mark rising/falling success rates. Bold/underlining mark best/second-best rates per family; blue shading marks R3.

| Error family       | Round | Attempted | Successful | Success (%)    |
|--------------------|-------|-----------|------------|----------------|
| Visual misreading  | R1    | 12,189    | 9,826      | <b>80.61</b>   |
|                    | R2    | 10,245 ↓  | 6,789 ↓    | 66.27 ↓        |
|                    | R3    | 10,478 ↑  | 7,094 ↑    | <u>67.70</u> ↑ |
| Reasoning error    | R1    | 9,355     | 5,279      | 56.43          |
|                    | R2    | 7,857 ↓   | 6,246 ↑    | <u>79.50</u> ↑ |
|                    | R3    | 6,789 ↓   | 5,647 ↓    | <b>83.18</b> ↑ |
| Differential error | R1    | 4,064     | 2,995      | <b>73.70</b>   |
|                    | R2    | 6,742 ↑   | 4,673 ↑    | <u>69.31</u> ↓ |
|                    | R3    | 4,786 ↓   | 3,291 ↓    | 68.76 ↓        |
| Knowledge gap      | R1    | 1,625     | 1,467      | <b>90.28</b>   |
|                    | R2    | 987 ↓     | 550 ↓      | 55.72 ↓        |
|                    | R3    | 1,567 ↑   | 765 ↑      | 48.82 ↓        |

## C Medical Taxonomy and Error Categories

### C.1 Taxonomy, Permissions, and Execution Contracts

#### C.1.1 Four-Dimensional Medical Tag Dictionary

The medical taxonomy describes each question along four independent axes: modality, anatomy, disease, and difficulty. Each axis allows multiple hierarchical labels for capability analysis and question retrieval.

*Modality hierarchy.* Computed tomography (CT), magnetic resonance imaging (MRI), positron emission tomography (PET), and single-photon emission computed tomography (SPECT) denote imaging modalities. Other abbreviations used below are optical coherence tomography (OCT), fundus fluorescein angiography (FFA), indocyanine green angiography (ICGA), cone-beam computed tomography (CBCT), electroencephalogram (EEG), and electromyogram (EMG). ECG denotes electrocardiogram, as in Fig. 3.

| Parent             | Second-level scope                                                         |
|--------------------|----------------------------------------------------------------------------|
| 1 Radiography      | X-ray; CT; MRI; mammography; fluoroscopy; angiography                      |
| 2 Nuclear medicine | PET/CT; SPECT; radionuclide scans                                          |
| 3 Ultrasound       | B-mode; echocardiography; vascular; interventional                         |
| 4 Endoscopy        | Gastrointestinal; bronchial; laryngeal/nasal; bladder; joint; laparoscopic |
| 5 Pathology        | Histology; cytology; molecular pathology                                   |
| 6 Ophthalmology    | Fundus photography; OCT; FFA/ICGA; slit lamp                               |
| 7 Dermatology      | Dermoscopy; clinical photography; reflectance confocal microscopy          |
| 8 Dental           | Panoramic; CBCT; periapical; cephalometric                                 |
| 9 Functional tests | ECG; EEG; EMG                                                              |

### *Anatomy hierarchy.*

| Parent                 | Second-level scope                                                                |
|------------------------|-----------------------------------------------------------------------------------|
| 1 Nervous system       | Brain; spinal cord; peripheral nerves                                             |
| 2 Head and neck        | Orbit; ear; sinuses; oral/maxillofacial; pharynx/larynx; neck soft tissue         |
| 3 Thorax               | Lung/airway; mediastinum; pleura; heart; great vessels; breast; chest wall        |
| 4 Abdomen              | Hepatobiliary; pancreas; spleen; gastrointestinal; renal/urinary; retroperitoneum |
| 5 Pelvis               | Male reproductive; female reproductive; pelvic floor; pelvic bones                |
| 6 Spine                | Cervical; thoracic; lumbar; sacrococcygeal; disc; canal; paraspinal tissue        |
| 7 Musculoskeletal      | Craniofacial bones; upper limb; lower limb; muscles; periarticular soft tissue    |
| 8 Vascular             | Arteries; veins; lymphatic system                                                 |
| 9 Skin and soft tissue | Head/face; trunk; limbs; subcutaneous; deep soft tissue                           |

### *Disease hierarchy.*

| Parent                       | Second-level scope                                                                                                                                                                               |
|------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1 Neoplasm                   | Malignant; benign; metastatic; precancerous categories                                                                                                                                           |
| 2 Infection and inflammation | Bacterial; tuberculosis; viral; fungal; parasitic; noninfectious inflammation                                                                                                                    |
| 3 Vascular disease           | Ischemic; hemorrhagic; aneurysm; thrombosis/embolism; malformations; atherosclerosis/stenosis                                                                                                    |
| 4 Degenerative disease       | Joint; neural; organ degeneration                                                                                                                                                                |
| 5 Trauma                     | Fracture; joint injury; visceral injury; head injury; foreign body                                                                                                                               |
| 6 Congenital/developmental   | Cardiovascular; nervous; skeletal; digestive; urinary; respiratory                                                                                                                               |
| 7 Metabolic/endocrine        | Diabetes complications; thyroid; adrenal; gout; metabolic bone                                                                                                                                   |
| 8 Autoimmune                 | Lupus; ankylosing spondylitis; multiple sclerosis; neuromyelitis optica spectrum disorder (NMOSD); autoimmune liver; immunoglobulin G4 (IgG4); vasculitis; systemic sclerosis; myositis; Sjogren |
| 9 Functional disorders       | Obstruction; effusion; ectopic gas; torsion; stones                                                                                                                                              |
| 10 Iatrogenic/postoperative  | Normal postoperative change; complications; implants; post-radiotherapy; drug-related                                                                                                            |

*Leaf specificity and multi-label assignment.* CT branches into noncontrast, contrast-enhanced, CT angiography (CTA), high-resolution CT (HRCT), and low-dose CT; MRI includes T1- and T2-weighted imaging, fluid-attenuated inversion recovery (FLAIR), diffusion-weighted imaging (DWI), magnetic resonance angiography (MRA), functional MRI (fMRI), magnetic resonance spectroscopy (MRS), and susceptibility-weighted imaging (SWI). Anatomy can descend from thorax to lung/airway and then to a lung lobe, or from abdomen to pancreas and then head/body/tail. Disease labels descend from a broad family to a mechanism or disease group and then to a named entry. Preserve image IDs for mixed modalities and record every supported anatomical site. **Assign the deepest supported node**, not a guessed leaf; use unknown, not-applicable, or a supported parent when necessary. **Expand ancestors as a set** and count each question once per group. Related concepts at different paths require an explicit alias ledger, not automatic merging. For each label, store source text, evidence span/image ID, confidence, visibility, and taxonomy version.

### *Difficulty rubric.*



| Level | Task requirement                                                                                      |
|-------|-------------------------------------------------------------------------------------------------------|
| L1    | Single modality and conspicuous canonical findings; direct recognition.                               |
| L2    | Common presentation requiring localization or characterization and limited context.                   |
| L3    | Atypical or multiple findings requiring sequence integration or differential comparison.              |
| L4    | Subtle findings, multisystem involvement, or complex context requiring extensive integration.         |
| L5    | Highly complex or disputed tasks requiring multidisciplinary, quantitative, or specialized knowledge. |

Assign difficulty from the actual question: visibility, number of required comparisons, evidence integration, and specialized knowledge. A disease name, modality, or a single failed rollout does not determine the level. Rare disease alone does not force L5. Record the rationale and adjudicate ambiguous adjacent levels.

### C.1.2 Error Codebook: Four Parents and Twenty Subtypes

We specify five subtypes under each of the four existing parent categories. Each row defines an observable diagnostic criterion and a navigation operation; it does not supply a gold answer.

*V: Visual misreading.*

| Code / subtype                                  | Diagnostic criterion                                                               | Navigation operation                                       |
|-------------------------------------------------|------------------------------------------------------------------------------------|------------------------------------------------------------|
| <b>V1 Missed finding</b>                        | A supported salient finding is not observed or described.                          | <b>Reinspect</b> the relevant region.                      |
| <b>V2 Hallucinated or misidentified finding</b> | An artifact, normal structure, or unsupported feature is read as a lesion.         | <b>Check</b> direct visual support.                        |
| <b>V3 Anatomy or laterality error</b>           | The organ, level, compartment, or side conflicts with visible evidence.            | <b>Recheck</b> landmarks and orientation.                  |
| <b>V4 Modality, sequence, or phase error</b>    | Acquisition type, sequence, or phase is misidentified, distorting feature reading. | <b>Verify</b> acquisition context before reading features. |
| <b>V5 Cross-image correspondence error</b>      | A lesion is matched incorrectly across images, slices, or views.                   | <b>Align</b> images and anatomical locations.              |

*R: Reasoning error.*

| Code / subtype                                    | Diagnostic criterion                                                                | Navigation operation                                    |
|---------------------------------------------------|-------------------------------------------------------------------------------------|---------------------------------------------------------|
| <b>R1 Omitted evidence</b>                        | An observed finding or stated fact is omitted from inference.                       | <b>Reintroduce</b> the omitted evidence.                |
| <b>R2 Conflicting evidence integration</b>        | Correctly extracted evidence conflicts and is not reconciled.                       | <b>Reconcile</b> supporting and opposing evidence.      |
| <b>R3 Unsupported inference or causal leap</b>    | A conclusion lacks the required logical or causal link.                             | <b>Restore</b> a testable intermediate inference.       |
| <b>R4 Temporal or quantitative error</b>          | Correctly read facts are combined with an incorrect temporal or numerical relation. | <b>Check</b> chronology, units, or calculation.         |
| <b>R5 Task constraint or answer-mapping error</b> | The reasoning violates the requested target, negation, or option mapping.           | <b>Recheck</b> the requested target and answer mapping. |

*D: Wrong differential.*

| Code / subtype                              | Diagnostic criterion                                                | Navigation operation                             |
|---------------------------------------------|---------------------------------------------------------------------|--------------------------------------------------|
| <b>D1 Omitted candidate</b>                 | A necessary evidence-compatible alternative is not considered.      | <b>Broaden</b> the candidate comparison.         |
| <b>D2 Missing discriminative comparison</b> | Alternatives are listed without comparing discriminative evidence.  | <b>Identify</b> discriminative evidence.         |
| <b>D3 Premature closure or anchoring</b>    | An initial diagnosis is fixed before counterevidence is considered. | <b>Revisit</b> counterevidence and alternatives. |
| <b>D4 Incorrect candidate ranking</b>       | Known support, counterevidence, or priors are weighted incorrectly. | <b>Compare</b> the strength of evidence.         |

| Code / subtype                 | Diagnostic criterion                                                                         | Navigation operation                         |
|--------------------------------|----------------------------------------------------------------------------------------------|----------------------------------------------|
| <b>D5 Misapplied exclusion</b> | An alternative is excluded using insufficient negative evidence or a nonessential criterion. | <b>Check</b> whether exclusion is justified. |

*K: Knowledge gap.*

| Code / subtype                                          | Diagnostic criterion                                                       | Navigation operation                                  |
|---------------------------------------------------------|----------------------------------------------------------------------------|-------------------------------------------------------|
| <b>K1 Missing finding interpretation</b>                | The required medical meaning of a finding or term is missing.              | <b>Cue</b> the relevant conceptual relation.          |
| <b>K2 Missing diagnostic criterion</b>                  | A required diagnostic, grading, or staging rule is unavailable.            | <b>Cue</b> the relevant rule category.                |
| <b>K3 Missing anatomy or pathophysiology</b>            | Knowledge linking anatomy, mechanism, and manifestations is missing.       | <b>Bridge</b> the minimal mechanistic relation.       |
| <b>K4 Missing applicability condition or exception</b>  | A general rule is known but its prerequisites or exceptions are missing.   | <b>Cue</b> the applicability check.                   |
| <b>K5 Missing investigation or management knowledge</b> | A next-investigation or management task requires an unavailable principle. | <b>Cue</b> the decision principle without the answer. |

*Boundary rules and reporting.* A finding not perceived is V1; a finding already perceived but omitted from inference is R1. Incorrect correspondence across images is V5; conflicting correctly extracted evidence is R2. Candidate-set construction, comparison, and exclusion belong to D; a demonstrated missing prerequisite rule belongs to K. Knowledge-gap labels require evidence of a missing prerequisite rule. **Select the earliest consequential supported error** as the primary subtype and record later failures as secondary labels, with step spans and evidence. If the parent is clear but the subtype is not, retain the parent and an unresolved-subtype status; if even the parent is unsupported, mark the audit unresolved. These are annotation states, not extra error classes. Use **one primary subtype per trajectory** for additive counts; report secondary labels separately. Aggregate V1–V5 into V, and similarly for R, D, and K. Re-audit fresh rollouts at each checkpoint using subtype-specific annotations and denominators.

### C.1.3 Information Access Matrix

| Component                            | Permitted inputs                                                                                               | Excluded inputs                                                                            |
|--------------------------------------|----------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------|
| <b>Feedback evaluator</b>            | Feedback items, labels, current predictions; emits aggregates and transitions                                  | Final-test results for planning; feedback export into SFT                                  |
| <b>Planner / retrieval</b>           | Aggregate feedback, training tag metadata (including true disease labels), coverage, budgets                   | Final-test items/results; any authority to forward answer-derived labels to repair         |
| <b>Auditor / hint renderer</b>       | Original images, question/options, current failed trajectory, entropy profile, observable tags with provenance | Gold answer, answer-derived disease/location, reference rationale, hidden retrieval labels |
| <b>Continuation <math>M_t</math></b> | Original input and assistant prefill $p_b \oplus g(h)$                                                         | Gold answer, full audit, hidden tags, teacher solution                                     |
| <b>Admission / exporter</b>          | Candidate, <b>reference answer</b> for verification, boundary, provenance, suffix scores                       | Feedback/test examples as training targets                                                 |
| <b>SFT trainer</b>                   | Frozen training manifest; $p_b \oplus g(h) \oplus s^+$ , masks                                                 | Full audit as model input; feedback/test samples                                           |

Use **explicit allowlist serialization**, not deletion of one answer field from an otherwise unrestricted object. **Reject fields without observable provenance**, including filenames/IDs encoding labels. An authorized evaluator can identify a wrong training answer before repair without transmitting the answer. Generation sees only a boolean routing decision, not the key. Record hashes of redacted requests for leakage audit. The planner’s chosen hint family is passed through an eligibility check derived from observable context and current diagnosis; disease-bearing planner explanations are never forwarded.

### C.1.4 Structured Fields and Plan Validation

The execution interface requires JavaScript Object Notation (JSON) records containing the following fields. An omitted measurement is null with a reason; zero means measured zero.

**Question metadata:** canonical question ID; source ID; patient/case/article cluster IDs when available; image hashes and near-duplicate cluster; split; taxonomy version; per-image modality; all four tag sets; per-tag provenance, visibility (statistical or observable), and uncertainty.

**Feedback record:** checkpoint hash; split hash; group and resolved parent; unique sample count; error count/rate; all four transition counts; both regression denominators; current error-type/subtype counts; confidence interval; training-pool unique availability; measured repair/admission/cost fields.

**Plan:** round ID; generator and initialization hashes; planner type/version and seed; feedback hash; target group IDs; numeric feedback evidence and fallback ancestry; training sources; selected IDs; attempt/call/token ceilings; per-group quota; permissible hint operation; history replay allocation; replacement rule; supervised-token cap; shortfall/reallocation policy; **fixed training-config hash**; acceptance-validator result.

**Repair record:** question and checkpoint IDs; rollout hash; step and token spans; entropy profile and proposal; current error type/subtype; audit boundary and exact prefix span; observable-tag payload hash; recovery cue; each continuation and seed; correctness and integrity checks; current-model suffix scores and cutoffs; admission/rejection reason; all generated/scored tokens.

**Mix and training record:** ordered example IDs; source round and generator; dedup/replacement ledger; historical/new role; token masks; sequence lengths; unit weights; packing map; effective repetitions and exposures; optimizer and scheduler reset flags; learning rate/warmup; initialization hash; output hash; second-stage and cumulative costs.

**Execution receipt:** planned and actual quotas, unique selected and attempted failures, successful and admitted repairs, substitutions, exhausted budgets, rejected records, uncovered groups, and actual training tokens.

**The validator rejects** cross-split IDs, duplicate quota charges, budgets above caps, out-of-pool retrieval, unsupported hint operations, missing visibility provenance, or changed training hyperparameters. Save the validated plan before generation and record actual execution counts afterward.

The complete feedback-agent prompt is given in Appendix D.2.

### C.1.5 Training Configuration and Evaluation Scope

Training uses brain floating-point 16 (BF16) precision on graphics processing units (GPUs). Parallelism abbreviations are tensor parallelism (TP), expert parallelism (EP), pipeline parallelism (PP), context parallelism (CP), expert tensor parallelism (ETP), and data parallelism (DP); the following number gives the parallelism degree. Configuration keys retain their literal software spelling; llm denotes large language model, vit denotes vision transformer, cpu denotes central processing unit, moe denotes mixture of experts, gemm denotes general matrix multiplication, mlp denotes multilayer perceptron, and lr denotes learning rate.

**Table 14** Recorded MoE training configuration.

| Configuration                            | Value                          | Configuration           | Value              |
|------------------------------------------|--------------------------------|-------------------------|--------------------|
| Framework                                | ms-swift + Megatron-LM backend | add_non_thinking_prefix | true               |
| Precision                                | BF16                           | loss_scale              | ignore_empty_think |
| micro_batch_size                         | 1                              | MAX_PIXELS              | 3145728            |
| global_batch_size                        | 64                             | dataset_num_proc        | 96                 |
| packing                                  | true                           | dataloader_num_workers  | 8                  |
| attention_backend                        | flash                          | logging_steps           | 1                  |
| calculate_per_token_loss                 | true                           | eval_iters              | 1000               |
| freeze_llm / freeze_vit / freeze_aligner | false / true / true            | recompute_granularity   | selective          |
| optimizer_cpu_offload                    | true                           | sequence_parallel       | true               |
| use_precision_aware_optimizer            | true                           | max_length              | 28672              |
| offload_fraction                         | 1.0                            | epochs                  | 2                  |

| Configuration             | Value | Configuration      | Value                            |
|---------------------------|-------|--------------------|----------------------------------|
| moe_grouped_gemm          | true  | save_steps         | 1000 (pilot: 100)                |
| moe_shared_expert_overlap | true  | save_total_limit   | 2                                |
| moe_aux_loss_coeff        | 1e-6  | Recompute modules  | core_attn + moe + mlp            |
| lr                        | 1e-5  | moe_permute_fusion | true                             |
| lr_warmup_fraction        | 0.05  | Parallelism        | TP4/EP8/PP1/CP1/ETP1;<br>64 GPUs |
| min_lr                    | 1e-7  |                    |                                  |

EP and MoE-specific options apply only to MoE models.

Both example types use full-response loss, retaining recovery cues in repairs.

§4 reports accuracy on seven benchmarks and their mean, comparing each backbone’s best complete checkpoint with Base.

## D Prompts

### D.1 Shared Question-Answering Prompt

All question-answering tasks use the following instruction verbatim.

#### Shared Question-Answering Instruction

Please reason step by step, and put your final answer within `\boxed{\}`.

### D.2 Main-Agent Feedback and Next-Round Planning

The following template connects feedback diagnosis to next-round selection and repair priorities under the method protocol. It specifies agent behavior; run-specific inputs and execution records are required to document its use.

#### Main Agent: Feedback Summary and Next-Round Data Plan

You are the feedback and data-planning agent for MedEvo. Diagnose the current model’s weaknesses and produce an executable next-round data plan. Do not write solutions to training questions or change the experimental comparison settings.

**Inputs.** Current round and checkpoint; the fixed 5,000-question feedback-set results for current and previous checkpoints; four-axis medical tags; four error families and their twenty subtypes; per-group sample counts, correctness transitions, and current error diagnoses; available training-question IDs and tags; historical trajectory records; fixed generation/repair budgets, support threshold, regression weight  $\alpha$ , and admission rules. Missing inputs must be reported, never invented.

**1. Summarize feedback.** For each supported tag group, report error rate, change from the previous round, corrected errors, regressions, and persistent errors. Summarize four-family and twenty-subtype distributions with explicit denominators. Count each question once within a group; do not sum overlapping groups into an overall rate. Merge unsupported groups into supported parents. Distinguish observed changes from explanations that require further evidence.

**2. Set priorities.** Rank groups using  $S_{g,t} = \epsilon_{g,t} + \alpha N_{g,t}^{10}/n_g$  under the supplied configuration. Identify dominant or persistent error types within each priority group. State which measured feedback supports each priority; do not infer that a high-error modality contains every error subtype.

**3. Select training questions.** Allocate the fixed budget proportionally to group scores, cap quotas by available unique questions, and redistribute unused quota. Preserve modality coverage and apply the configured fallback when all scores are zero or support is insufficient. Retrieve by medical tags, using historical errors only as prioritization cues. Deduplicate multi-tag questions and return selected training IDs, group quotas, and any shortfall. **Never select feedback or final-test examples.**

**4. Guide fresh rollouts and repair.** Ask the current checkpoint to generate the configured eight rollouts per selected question. **Diagnose their actual errors again.** Use feedback-identified modality–error pairs to prioritize failed training trajectories: prefer joint matches, followed by modality-only or error-type-only matches

when necessary. For each selected trajectory, pass its observable medical tags and current error diagnosis to the hint generator. The recovery hint must address the error actually observed in this trajectory. Send observable tags and local diagnoses to the auditor and hint generator, **without reference answers or answer-bearing metadata**. Retain correct rollouts and accepted continuations under the existing screening rules; if repair fails, use an available quality-screened correct rollout. If none is available, retain the valid historical target or report no accepted target.

**5. Plan the next round.** **Keep at most one active target per question** and replace it only with an accepted improvement. **Train on the complete updated dataset from  $M_0$ .** Keep training settings and evaluation protocol fixed. Evaluate the next checkpoint on the same feedback set to measure the targeted changes; independent benchmark results are reported separately and must not drive selection.

**Output.** Return a structured plan containing: round and checkpoint; feedback summary and denominators; priority tag/error pairs and supporting statistics; score and quota per group; selected unique training IDs; rollout and repair budgets; repair priorities; dataset-retention/replacement rules; initialization  $M_0$ ; next-round feedback comparisons; and missing inputs or budget shortfalls. After execution, record actual generated, repaired, admitted, and replaced counts.

### D.3 Stage 1 Audit Prompt

The diagnostic model receives the image, the question with options, the failed trajectory, the complete step-level uncertainty profile, and the entropy-derived anchor and candidate window. The auditor is gold-blind: **the reference answer is not an input to diagnosis**. Answer-derived disease labels, hidden lesion positions, reference explanations, and planner retrieval metadata are also excluded. Only provenance-checked observable tags are allowed. The entropy proposal is a localization prior, not an error label: the auditor inspects the proposed region first, traces its evidence dependencies backward, and may move the final boundary when the medical or logical evidence requires it. It returns a structured audit that identifies a compact revision region and constrains the subsequent continuation. The emitted hint must remain answer-blind. Placeholders are shown in braces and are filled at runtime.

#### Prompt D.3: Stage-1 Entropy-Guided Medical Trajectory Audit

You are a senior medical reviewer auditing a failed multimodal reasoning trajectory segmented into semantic steps  $s_1, \dots, s_K$ .

You are given:

- (1) the medical image(s);
- (2) the clinical question and answer options;
- (3) the failed student trajectory segmented into semantic steps;
- (4) STEP\_ENTROPY: the uncertainty score  $E_k^*$  of every step;
- (5) NORMALIZED\_ENTROPY: the normalized uncertainty score of every step;
- (6) LOCAL\_SURGE: the local increase  $\Delta E_k$  relative to the preceding step history;
- (7) ENTROPY\_ANCHOR: the step proposed by the uncertainty locator;
- (8) ENTROPY\_WINDOW: the compact candidate window surrounding that anchor, and ENTROPY\_THRESHOLD if a threshold was used.
- (9) OBSERVABLE\_TAGS: per-image modality and other context supported by the original input, with evidence provenance; never hidden answer-derived tags.

The entropy signal guides audit localization. High uncertainty can reflect a genuine failure, but it can also arise from difficult terminology, ambiguous visual evidence, or a medically valid differential diagnosis. Combine the entropy proposal with the image, clinical evidence, and dependencies between reasoning steps to identify the first consequential failure.

**Error-conditioned hint construction.** Construct the hint from OBSERVABLE\_TAGS, ERROR\_TYPE, ERROR\_SUBTYPE, and KEY\_ERROR. Select a recovery operation addressing the diagnosed failure:



- **Visual misreading:** direct attention to the relevant region, visual attribute, or correspondence across available images.
- **Reasoning error:** revisit an unsupported inference or reconnect observed evidence to the requested decision.
- **Differential diagnosis error:** compare discriminative evidence or reconsider an insufficiently examined alternative.
- **Knowledge gap:** identify the relevant concept, rule, or applicability condition to reconsider.

Specify one concrete next operation grounded in the current question and trajectory. Express it as a brief recovery cue without supplying the corrected finding, final diagnosis, or answer option.

Audit procedure:

- (1) Inspect ENTROPY\_WINDOW first. Determine whether its local entropy peak or surge corresponds to an unsupported observation, invalid inference, incorrect differential comparison, or missing medical principle.
- (2) Trace evidence dependencies backward from ENTROPY\_ANCHOR. Decide whether the decisive failure begins inside the proposed window or an earlier step introduced the unsupported premise that caused the later entropy increase.
- (3) Do not automatically select the highest-entropy step. Justify the final AUDIT\_ANCHOR and SAFE\_PREFIX\_END using the medical image, clinical information, and logical dependencies.
- (4) Prefer a compact localized repair and preserve every preceding step that remains medically and logically valid.
- (5) You may move the audit anchor outside ENTROPY\_WINDOW when the proposal misses the first consequential failure, but briefly explain why the entropy signal was misleading.

Return a JavaScript Object Notation (JSON) object with the following fields, grounded strictly in the provided image and text:

- (a) ACTION: REPAIR or SKIP.
- (b) STEP\_SUMMARY: a brief one-sentence semantic summary for each step.
- (c) ENTROPY\_ASSESSMENT, containing: proposed\_anchor, proposed\_window, proposal\_status, and brief\_reason. proposal\_status must be one of {HIT, PARTIAL\_HIT, MISS, NON\_ERROR\_UNCERTAINTY}. HIT means that the first consequential failure is at or immediately adjacent to the proposed anchor; PARTIAL\_HIT means that the proposal covers the propagated error region but the unsupported premise begins earlier; MISS means that the first failure lies outside the proposed window; and NON\_ERROR\_UNCERTAINTY means that the high-entropy region is medically or logically valid.
- (d) AUDIT\_ANCHOR: the index of the first consequentially incorrect step.
- (e) SAFE\_PREFIX\_END, containing step\_index and the exact verbatim\_ending\_span of the last fully valid step before the repair region.
- (f) AUDIT\_WINDOW: a short contiguous sequence of steps that should be reconsidered.
- (g) ERROR\_TYPE and ERROR\_SUBTYPE: checkpoint-specific diagnosis, plus KEY\_ERROR: the localized error stated without revealing the correct answer.
- (h) HINT\_TYPE: one of {perceptual\_reexamination, inference\_reconstruction, differential\_disambiguation, knowledge\_bridging}.
- (i) MINIMAL\_HINT: a brief recovery cue with flexible wording, conditioned on observable medical tags and the diagnosed local error. Name the evidence or reasoning operation to revisit. Do not include the error-category label or audit explanation. Introduce the cue with a natural transition; no fixed phrase or first-person form is required.
- (j) IMAGE\_CHECK, containing applicable (true or false) and target (the specific modality, anatomy, laterality, finding, relation, or measurement to re-inspect).
- (k) KNOWLEDGE\_SCOPE: the narrow medical principle to recall, or null.
- (l) DIFFERENTIAL\_FOCUS: the discriminative axis to compare, or null.
- (m) STOP\_CONDITION: the evidence condition under which the continuation should commit to an answer and stop.

**Constraints:**

- Do not access hidden disease labels or hidden lesion coordinates.
- Do not state a gold answer, corrected diagnosis, or correct option letter.
- Do not provide a replacement chain of thought or teacher-authored solution.
- Do not introduce symptoms, history, tests, measurements, or visual findings absent from the input.
- Do not discard a valid step merely because its entropy is high, and do not preserve an invalid step merely because its entropy is low.
- SAFE\_PREFIX\_END must precede the first consequential failure.
- AUDIT\_WINDOW should normally overlap ENTROPY\_WINDOW, but may deviate when supported by explicit medical or logical evidence.
- If only final-answer formatting or option selection is wrong and no reasoning revision is justified, set ACTION to SKIP.

**D.4 Stage 2 Continuation Prompt**

The continuation input contains the original question and an unfinished assistant message with the retained prefix and an error-conditioned recovery cue. Stage 1 constructs the cue from observable medical tags, ERROR\_TYPE, ERROR\_SUBTYPE, and KEY\_ERROR. Feedback determines repair priority; the selected trajectory’s own diagnosis determines cue content. The same student continues from this prefilled message.

**Prompt D.4: Stage-2 Assistant-Prefill Continuation****User message:**

{original medical image(s), question, and answer options}

**Assistant message (unfinished):**

{safe\_prefix}

{Natural transition} {MINIMAL\_HINT from Stage 1}

MINIMAL\_HINT is inserted verbatim after the retained prefix. The student receives only the operational cue; structured error labels and the audit remain metadata. Generation continues immediately after the cue within the same assistant message, without repeating or modifying the prefix.

The SFT export retains the student-style recovery cue. Repaired and original correct examples use full-response loss in Equation 13, supervising all assistant-response tokens. No audit or external hint is supplied at final evaluation. Record exported examples and token masks with the training configuration.

**E Case Study**

We present two medical image question-answering cases, including the question, options, and student model output, with the hint marked within the complete reasoning text. We preserve the wording of the student outputs without language editing. Purple labels separate the original text, hint, and continuation; blue emphasis marks selected observations and reasoning actions in the quoted outputs.

**E.1 Case 08: Chest radiograph**


### Original question and options

A 67-year-old man is recovering in the surgical intensive care unit following an esophagectomy. A portable anteroposterior chest radiograph (Figure A) is obtained to evaluate his postoperative status. The patient has a right-sided chest tube in place. Based on the radiographic findings, which of the following is the most likely cause of the bilateral lower lung opacities?

Options:

- A. Pleural effusion
- B. Pulmonary embolism
- C. Emphysema
- D. Cavitory tuberculosis
- E. Tension pneumothorax

Please reason step by step, and put your final answer within \boxed{ }.



### Student model output

#### Original context and initial reading

<think>

The user wants a step-by-step reasoning and a final boxed answer. Let me look at the radiograph carefully. It's a **portable AP chest radiograph** of a 67-year-old man post-esophagectomy. There's a right-sided chest tube. There are ECG leads overlying the chest. There's an NG tube or feeding tube coursing down the midline (post-esophagectomy, possibly gastric pull-up). The question asks about **bilateral lower lung opacities**. Looking at the image, the lower zones do not show convincing increased opacity; the costophrenic angles are not clearly blunted, and there is no definite meniscus-like contour. The hemidiaphragms remain relatively well defined. So pleural effusion is not strongly supported.

#### Navigation hint

Before weighing the choices, **re-check the morphology of the basal opacities: costophrenic-angle blunting, a meniscus-like contour**, and partial obscuration of the hemidiaphragms would localize the process primarily to the...

#### Student continuation after the hint

Now, **let me think about the other options**: pulmonary embolism doesn't usually cause bilateral lower lobe opacities (may cause atelectasis or infarction, but not typical). Emphysema causes hyperlucency, not opacity. Cavitory TB causes upper lobe cavities, not bilateral lower lobe opacities. Tension pneumothorax causes hyperlucency and mediastinal shift, not opacity. So the answer is A, pleural effusion.

The question says "bilateral lower lung opacities" — in a postoperative patient with a chest tube, bilateral pleural effusions (or a chylothorax after esophagectomy) are common. The radiographic appearance of basal opacities with meniscus and blunting of the costophrenic angles supports pleural effusion.

Therefore the

**Let me describe the radiographic findings in more detail**: bilateral lower zone opacities with a meniscus sign, blunting of the costophrenic angles, obscuration of the hemidiaphragms; the chest tube is on the right; the left side shows a similar opacity. No evidence of hyperlucency or pneumothorax; no cavitory lesions; no findings of emphysema (which would be lucent). PE is not a radiographic diagnosis on chest X-ray anyway and wouldn't present as bilateral basal opacities.

Therefore the answer is A.

I should present step-by-step reasoning: identify the findings, exclude the options, and conclude with A.

**Let me double-check the image description once more**: the left costophrenic angle is blunted with a meniscus; the right base shows opacity consistent with effusion and possibly some atelectasis. The question asks for the "most likely cause" — pleural effusion.

Also, after esophagectomy, patients often develop pleural effusions due to chylothorax or transudate from fluid shifts. So A is clinically plausible as well.

Format: the answer should be moderately detailed, with step-by-step reasoning, and a final answer. Use headings or bullet points sparingly. Let me write a concise structured analysis: findings, differential exclusion, conclusion. Then give the final  
Let me write it like this:

**Step 1 — Systematic reading:** portable AP film; cardiac silhouette is difficult to assess; bilateral lower zone opacities; blunting of the costophrenic angles with a meniscus-like upper border; obscuration of the hemidiaphragms; right-sided chest tube (likely for drainage of pleural fluid/air); no evidence of pneumothorax (no visceral pleural line with hyperlucency); no cavitary lesions; lung volumes are not hyperinflated.

**Step 2 — Match with options:**

- Pleural effusion: fits — dependent (basal) opacities with meniscus, blunting of the costophrenic angles, common after esophagectomy (including chylothorax).
- PE: not visible on CXR; wouldn't cause diffuse basal opacities.
- Emphysema: causes hyperlucency, flattened diaphragms — opposite of opacity.
- Cavitary TB: upper lobe cavitary lesions, not bilateral basal opacities.
- Tension pneumothorax: hyperlucent hemithorax, mediastinal shift away — not present.

**Conclusion: A.**

I might also mention that the chest tube on the right is for drainage of pleural fluid/air, and that the left effusion is likely also present.

Let me double-check the arrow in the image: it points to the left shoulder/clavicle area — probably just a marker. Not relevant.

Final

I'll write the response with clear reasoning. Format: short paragraphs or bullet points. Let me use a structured but concise approach</think>

**Most appropriate answer:** \boxed{A}.

## E.2 Case 31: Elbow radiograph

### Original question and options

A 32-year-old patient presents to the emergency department with acute right elbow pain, swelling, and limited range of motion following a fall onto an out-stretched hand. An anteroposterior radiograph of the right elbow is obtained (Figure A). What is the most likely diagnosis?

Options:

- A. Normal elbow radiograph
- B. Isolated elbow dislocation
- C. Elbow fracture
- D. Advanced osteoarthritis
- E. Osteochondritis dissecans

Please reason step by step, and put your final answer within \boxed{ }.



### Student model output

#### Original context and initial reading

<think>

The user wants a step-by-step reasoning for a radiograph interpretation question. **Let me carefully look at the image.** The image is labeled "R" (right) and shows an elbow radiograph. The question states it's an AP radiograph, though the image looks somewhat like an oblique or lateral-ish view, but let's focus on the pathology. Looking at the radiograph: I see the **distal humerus, the proximal radius, and proximal ulna.**

#### Navigation hint

Before committing, **trace the radial head and neck cortices** to verify that the horizontal lucency causes **true cortical interruption or fragment displacement** rather than projectional overlap, then **assess whether the radiocapitellar and ulnohumeral relationships** are...

### Student continuation after the hint

The radiocapitellar alignment appears normal (the radial head is opposite the capitellum), and the ulnohumeral articulation is also maintained, suggesting no dislocation. In case of an isolated dislocation, the radial head would be displaced from the capitellum. Here the alignment looks preserved, so an isolated dislocation is unlikely. The presence of a fracture line at the radial head suggests a radial head fracture, so the answer is C (elbow fracture). Osteoarthritis wouldn't present acutely after trauma, and there are no osteophytes. Osteochondritis dissecans is a chronic condition with a loose body, not an acute trauma presentation.

Therefore, the most likely diagnosis is C: elbow fracture (radial head fracture). Let's describe the reasoning steps:

**1. Verify patient history:** 32-year-old, fall on an outstretched hand (FOOSH) -> typical mechanism for radial head/neck fractures, coronoid fractures, lateral collateral ligament injuries, etc.

**2. Radiographic findings:** AP view of the right elbow shows a lucent line with cortical disruption at the radial head/neck, with some displacement/fragmentation - consistent with a radial head fracture. The radiocapitellar and ulnohumeral joints remain congruent, excluding dislocation. No evidence of advanced osteoarthritis (no joint space narrowing, osteophytes, subchondral sclerosis). No chronic loose body or cartilage defect suggestive of osteochondritis dissecans.

**3. Therefore the answer is C.**

Also, it's worth mentioning that the "fat pad sign" may be present (elevated anterior fat pad), indicating an intra-articular fracture/hemarthrosis, supporting the diagnosis of a fracture. In the image, there may be a faint anterior fat pad sign. But the key is the radial head fracture.

Let me double-check the image description once more: the image shows the humerus, ulna, and radius. In the proximal radius region, there's a clear horizontal white line - a fracture of the radial head. The radial head may also be comminuted. So the answer is C.

Let me also consider the possibility of a coronoid fracture. Regardless, it's an elbow fracture.

Thus, the final

Format: The answer should be moderately detailed, including step-by-step reasoning: mechanism, radiographic findings, exclusion of other options, conclusion. Use headings or bullet points. Final

Let me carefully describe the image: In the image, the radial head shows a transverse lucency consistent with a fracture; the radial head fragment is slightly displaced; the radiocapitellar relationship is preserved; the ulnohumeral angle is normal; no dislocation.

Let me also mention that the lateral fat pad sign supports an intra-articular fracture.

Let me write the response.

Also, the question states "anteroposterior radiograph," but the image is somewhat lateral; I won't overemphasize the view. Final

**\*\*Step 1 – Mechanism of injury:\*\*** A fall onto an outstretched hand (FOOSH) is the classic mechanism for a radial head/neck fracture, coronoid fracture, or lateral collateral ligament injury in a young adult.

**\*\*Step 2 – Systematic review of the radiograph:\*\***

- **\*\*Joint alignment:\*\*** The radiocapitellar and ulnohumeral relationships are preserved – the radial head remains opposite the capitellum and the ulna articulates normally with the trochlea. This argues **\*\*against an isolated elbow dislocation\*\*** (option B), which would show obvious displacement of these relationships.

- **\*\*Bony detail:\*\*** There is a transverse lucent line with cortical disruption and a small displaced fragment at the radial head/neck region – i.e., a **\*\*fracture of the radial head\*\***. A faint anterior fat-pad sign (sail sign) is also suggestive of an intra-articular fracture/hemarthrosis.

- **\*\*No degenerative changes:\*\*** There is no joint-space narrowing, osteophyte formation, or subchondral sclerosis, excluding **\*\*advanced osteoarthritis\*\*** (option D), which would also be unlikely to present as an acute traumatic event.

- **\*\*No chronic lesion:\*\*** There is no well-circumscribed subchondral defect with a loose body, which would suggest **\*\*osteochondritis dissecans\*\*** (option E); that condition is chronic/overuse-related, not acute post-traumatic.

- The film is clearly **\*\*not normal\*\*** (option A) given the visible fracture line and fragment.

**\*\*Step 3 – Synthesis:\*\*** Acute trauma + FOOSH mechanism + visible fracture at the radial head with preserved joint alignment = a fracture of the elbow (radial head fracture), not a dislocation or degenerative/chronic lesion.

**\*\*Most likely diagnosis: C. Elbow fracture\*\***</think>

**Most appropriate answer:** \boxed{C}.



## F Additional Related Work

This section expands the concise main-text review with background on medical evaluation, supervision quality, localized correction, and alternative forms of self-evolution. It also covers complementary work in agent safety, robotics, and data generation.

### F.1 Medical Reasoning, Distillation, and Data Selection

**Medical vision-language models** use biomedical instruction tuning [2, 3], generalist reasoning with **Lingshu** [11], and English–Arabic medical image understanding in **BiMediX2** [12]. Verifier-guided search, knowledge-enhanced synthesis, and reinforcement learning[36, 37, 38], together with parallel reasoning and multi-agent collaboration[39, 40, 41], improve reasoning capabilities. Process-level diagnostics and distillation audits distinguish answer correctness from reliable visual recognition, knowledge application, and evidence integration[14, 15, 42].

**Medical hallucinations** also expose limitations of evaluation. In controlled error-injection experiments, Farrow et al. [16] show that clinically relevant hallucinations can leave rubric scores unchanged when the errors fall outside predefined criteria. This finding motivates explicit factuality checks alongside rubric-based assessment, to assess clinically relevant errors beyond the rubric criteria.

**Process supervision** provides intermediate-step feedback[17], while trajectory synthesis and selection address rationale quality, difficulty, and distribution[43, 44]. **ChainPrune** [18] organizes self-generated reasoning paths into a tree and selects paths that reduce redundancy while preserving essential reasoning steps. Its distinction between token length and structural efficiency highlights the importance of assessing reasoning content when selecting trajectories. **On-policy distillation** emphasizes compatibility between student states and teacher capabilities[9]; reference-induced shortcuts, reasoning paths, and state coverage also shape supervision value[45, 46, 47]. **MedEvo** connects medical error diagnosis with student-centered supervision, using minimal teacher hints for local correction and retaining student reasoning for **self-distillation**.

### F.2 Learning from Failed Trajectories and Localized Repair

Failed trajectories contain reusable reasoning as well as errors. Self-Refine and Reflexion improve subsequent attempts through self-feedback and verbal memory without weight updates [48, 49]. **Learning from Mistakes (LEMA)** constructs mistake–correction pairs, **TrajFusion** combines incorrect and correct trajectories, and Exploring Expert Failures (EEF) extracts useful actions from failed demonstrations [6, 10, 50]. Hindsight Hint Distillation (**HHD**) and FailForge guide recovery with concise hints or skills instead of full teacher solutions [1, 51].

Local recovery also depends on where guidance enters and how it becomes supervision. **Error-localization** studies, backtracking-based distillation, Test-Time Scaling via Error Localization (**TTEL**), and Selective Regenerative Decoding investigate localization, backtracking corrections, and prefix reuse [5, 19, 20, 52]. Self-Correction via Reinforcement Learning (**SCoRe**) learns self-correction through reinforcement learning; Project Aletheia and Pyli-gent derive supervision from verified backtracking traces or search trees; micro-reflective self-distillation combines micro-reflective trajectories with capability-dependent selection [21, 53, 54, 55]. EvoThink and Reflective On-Policy Self-Distillation (ROSD) instead use preference learning and reflective on-policy self-distillation [56, 57].

Within this trajectory-learning tradition, **MedEvo** focuses on medical-contextualized, **same-student local navigation**: preserving supported prefixes, guiding continuation with gold-blind auditing and tag-conditioned hints, and selecting supervision by student-relative learnability.

### F.3 Data Evolution and Recursive Self-Improvement

Self-Taught Reasoner (**STaR**) and Reinforced Self-Training with Expectation–Maximization (ReST<sup>EM</sup>) iteratively construct supervision through self-generated reasoning and outcome selection[7, 8]; EvoThink and **ROSD** further explore preference learning and reflective self-distillation[56, 57]. **Capability-dependent selection** further connects the supervision distribution to what a student can learn [21].

**Recursive self-improvement** examines how feedback changes subsequent learning experience [58]. In clinical interaction, **MediSkill-Evo** [22] evolves validated clinical, process, symbolic, and visual knowledge banks without fine-tuning the backbone, using process constraints to guide evidence-grounded actions. **MedEvo** instead updates model weights through supervised fine-tuning on repaired student trajectories: it diagnoses medical groups and error patterns on a fixed feedback set, then generates and repairs training-pool trajectories for the next round of **self-distillation**.

## F.4 Failure Feedback beyond Medical Reasoning

Failure-aware approaches also organize recovery using structured failure feedback in agent-safety and robotics settings [59, 60, 61].

## F.5 Analysis of Frontier-Model Reasoning

Luo et al. [62] characterize externalized reasoning in frontier models through token efficiency, reasoning-step types, and reasoning-tree structure, complementing outcome-based evaluation with trajectory-level analysis.

## F.6 Broader Data and Environment Evolution

Difficulty-guided chart synthesis and multimodal generation workflows expand the supervision distribution [63, 64]. Data-strategy optimization, evolving environments, and autonomous medical curricula provide broader settings for recursive improvement [65, 66, 67]. These studies provide context for data evolution; MedEvo focuses on repairing student medical-reasoning trajectories and selecting subsequent supervision.

## G Additional Baseline Details

This section explains the roles of the comparison models and training baselines in Table 2. The main text reports benchmark results; the descriptions below distinguish external model comparisons from the trajectory-supervision baselines used to evaluate MedEvo.

### G.1 Backbones and External Comparison Models

The primary backbone is Qwen3.5-9B, with Qwen3.5-35B and Qwen3.5-397B used for the scale experiments [24]. A Base result refers to the corresponding initial checkpoint before MedEvo training. Later-round models are trained from that **same initial checkpoint** using the updated supervision dataset; the current checkpoint generates the trajectories used to construct that dataset.

The external comparison set includes general-purpose and medical-specialized models. Medical models include LLaVA-Med [2], HuatuoGPT-Vision [32], BiMediX2 [12], Lingshu [11], and HealthGPT [35]. InternVL2.5 and InternVL3 provide general-purpose multimodal comparisons [33, 34]. Table 2 groups models by proprietary access and open-source parameter scale. These model comparisons provide performance context; they are distinct from **controlled ablations** of the MedEvo training procedure.

### G.2 Trajectory-Supervision Baselines

*Baseline 1: HHD.* Hindsight Hint Distillation uses concise hints to scaffold recovery [1]. Its role in the comparison is to assess hint-based supervision alongside MedEvo’s **medical diagnosis, prefix preservation, and feedback-driven data updates**. The cited work supplies the underlying method; the medical-benchmark scores in Table 2 are those reported in this study.

*Baseline 2: TrajFusion.* TrajFusion combines incorrect and correct trajectories to construct training supervision [10]. It provides a comparison for learning from failed reasoning. MedEvo instead constructs a **corrected suffix** from the same student’s retained prefix and uses medical feedback to prioritize later repairs.

*Baseline 3: Teacher SFT.* This baseline uses supervised fine-tuning on supervision generated by Qwen-3.8-Max. It compares **teacher-generated supervision** with MedEvo’s repaired student trajectories. We construct this baseline for the supervision-source comparison.

### G.3 Reading the Comparisons

The medical icon identifies medical-specialized models, and the repair icon identifies trajectory-repair methods. Blue shading identifies MedEvo. Bold and underlining indicate the highest and second-highest distinct scores within each model group and column, with ties sharing a rank. The **seven-benchmark average** weights each benchmark equally. It is distinct from repair success, which measures successful corrections among questions submitted for repair.

Component ablations test full rewriting, suffix continuation, medical-tag feedback, and error-type feedback. Their purpose is to isolate design choices within the reported setup. Training configuration and information-access rules are provided in Appendices C.1.5 and C.1.4, respectively.